

DATA STRUCTURES

Internal Meta-data

Core Data Storage

Temporary Data Structures

Table Indexes

DESIGN DECISIONS

Data Organization

- How we layout data structure in memory/pages and what information to store to support efficient access.

Concurrency

- How to enable multiple threads to access the data structure at the same time without causing problems.

Concurrency means more than 1 thread is able to access data structure at some time without any problem.

Unordered

HASH TABLES (hash map)

A hash table implements an unordered
associative array that maps keys to values.

keys
(values) are mapped to
values.

It uses a hash function to compute an offset into
this array for a given key, from which the desired
value can be found.

Space Complexity: **O(n)** → To store the n elements we need
array.

Time Complexity:

- Average: **O(1)** ← *Databases need to care about constants!*
- Worst: **O(n)**

↳ Due to collision if all no.s point to some offset.

We use
hash
function to
find the
key or
offset.

STATIC HASH TABLE

→ Static

Size doesn't

change

∴ N is
fixed.

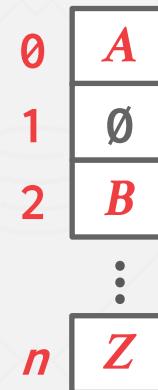
Allocate a giant array that has one slot for every element you need to store.

To find an entry, mod the key by the number of elements to find the offset in the array.

$\text{key} = \text{hash}(\text{value}) \% N \rightarrow \text{hash fn used.}$

key

$\text{hash(key)} \% N$

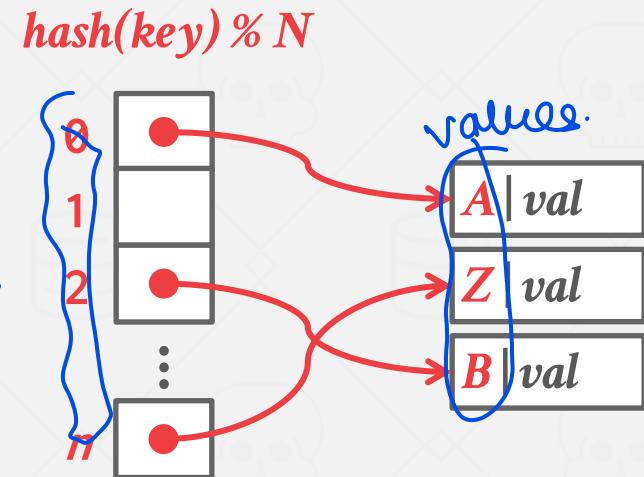


∴ static N is fixed we
allocate a giant array to
store element.

STATIC HASH TABLE

Allocate a giant array that has one slot for every element you need to store.

To find an entry, mod the key by the number of elements to find the offset in the array.



→ fixed size

→ unique
elements

ASSUMPTIONS

Assumption #1: Number of elements is known ahead of time and fixed.

← (Static definition itself)

Assumption #2: Each key is unique.

(No collision) ↴

Assumption #3: Perfect hash function.

→ If $\text{key}_1 \neq \text{key}_2$, then

$\text{hash}(\text{key}_1) \neq \text{hash}(\text{key}_2)$

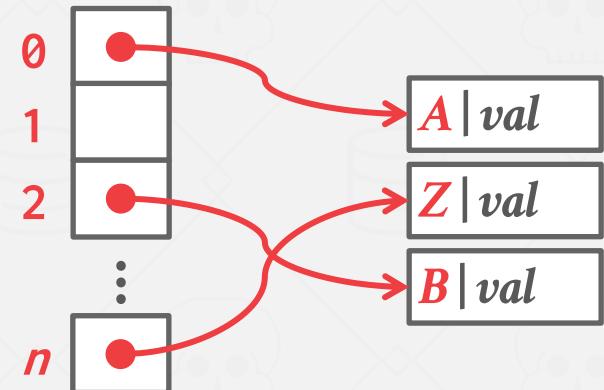
∴ Perfect
Hash Fxn.

i.e. $5 \neq 10$ but if $N=5$

$$\text{hash}(5) = 5 \rightarrow 5 = 0$$

∴ not perfect hash fn here.

$$= \text{hash}(10 \times 5)$$



$\text{hash}(\text{key}) \% N$

while designing

hashable

properly design

hash fn

properly
design

hashing
scheme

→ how to handle collisions.

HASH TABLE



Design Decision #1: Hash Function

- How to map a large key space into a smaller domain.
- Trade-off between being fast vs. collision rate.

in hashing scheme
tradeoff is

allocating large hash
table

(vs)

additional info to get
or put keys.

Design Decision #2: Hashing Scheme

- How to handle key collisions after hashing.
- Trade-off between allocating a large hash table vs. additional instructions to get/put keys.

in Hash fn the trade off is.

speed vs collision
rate.

TODAY'S AGENDA

Hash Functions

Static Hashing Schemes

Dynamic Hashing Schemes

For DBMS hash tables

we just use normal

HASH FUNCTIONS

hashing usage of cryptographic hashing will make it more complex.

For any input key, return an integer representation of that key.

↳ offset has to be returned for the input key.

We do not want to use a cryptographic hash function for DBMS hash tables (e.g., SHA-2).

We want something that is fast and has a low collision rate.

Hence we don't prefer cryptographic hashing

HASH FUNCTIONS

CRC-64 (1975)

- Used in networking for error detection.

MurmurHash (2008)

- Designed as a fast, general-purpose hash function.

Google CityHash (2011)

- Designed to be faster for short keys (<64 bytes).

Facebook XXHash (2012)

- From the creator of zstd compression.

← State-of-the-art

Google FarmHash (2014)

- Newer version of CityHash with better collision rates.

STATIC HASHING SCHEMES

Approach #1: Linear Probe Hashing

Approach #2: Robin Hood Hashing

Approach #3: Cuckoo Hashing

If collision happens we will linearly search for the next free slot in table.

LINEAR PROBE HASHING

Single giant table of slots. (similar to static Hashing).

Resolve collisions by linearly searching for the next free slot in the table.

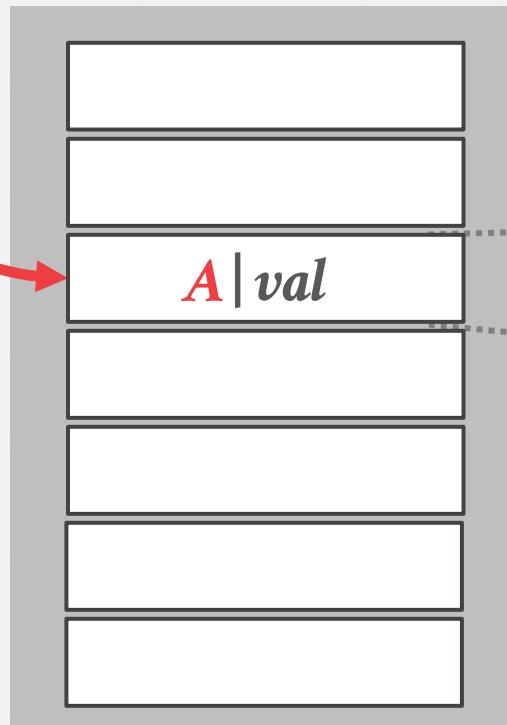
- To determine whether an element is present, hash to a location in the index and scan for it.
- Must store the key in the index to know when to stop scanning.
- Insertions and deletions are generalizations of lookups.

Search(value) ? ~~Offset = hash(value)~~ go to that offset ? Scan for it.

LINEAR PROBE HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F

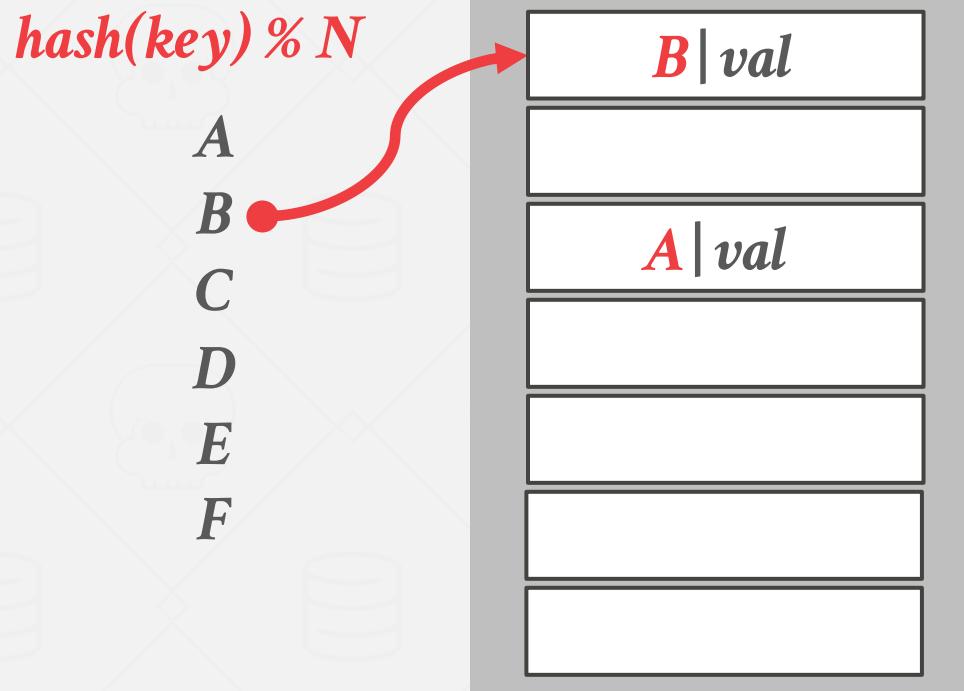


we store it as a pair

`<key|value>`

`<key> | <value>`

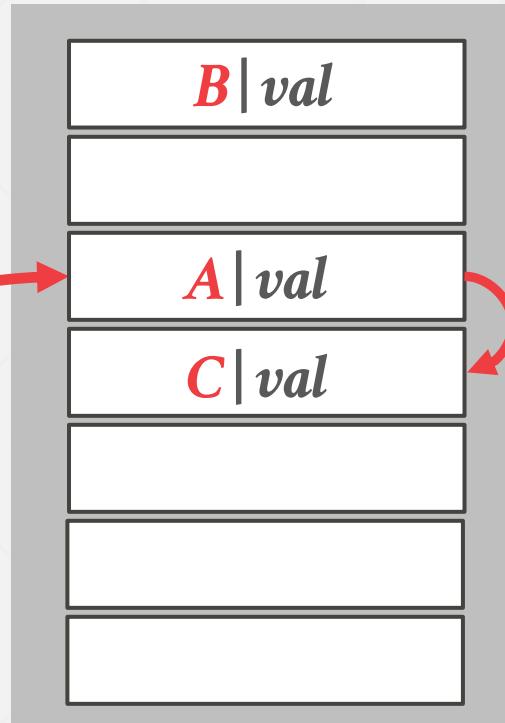
LINEAR PROBE HASHING



LINEAR PROBE HASHING

$\text{hash}(\text{key}) \% N$

A
 B
 C ●
 D
 E
 F



LINEAR PROBE HASHING

$\text{hash}(\text{key}) \% N$

In Search

Let us search

for x it is

at $\text{key } i$

now we go to i
as per hash fn

or per hash fn

but it is not there

A

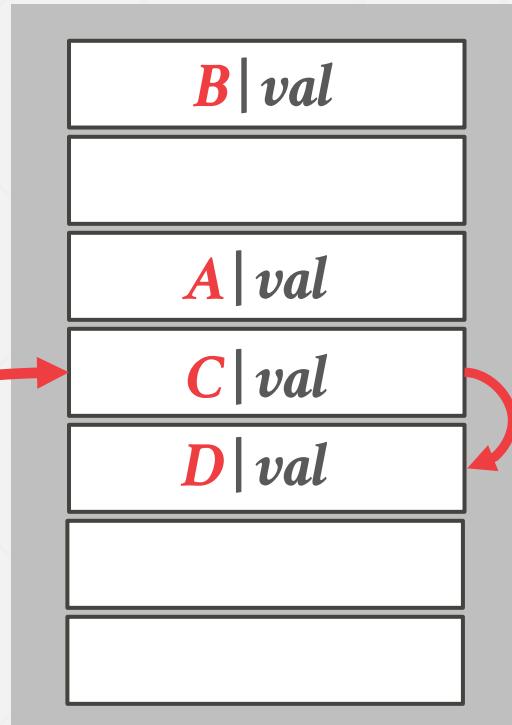
B

C

D

E

F



we search whole table.

We didn't find x how

to stop?

- if the index = i

then it means we

reached start again.

- Stop in middle

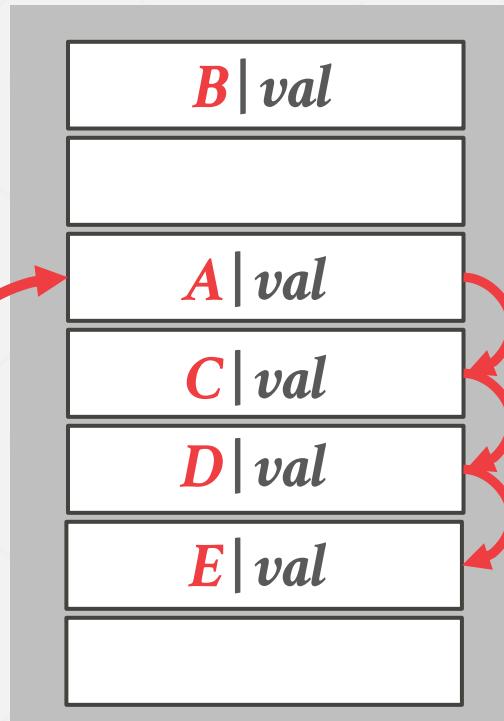
If we find x we

stop our search
there.

LINEAR PROBE HASHING

$\text{hash}(key) \% N$

A
B
C
D
E
F



See we are searching
for next available
free space in case
of collision

LINEAR PROBE HASHING

$\text{hash}(\text{key}) \% N$

In lookup
or

Search

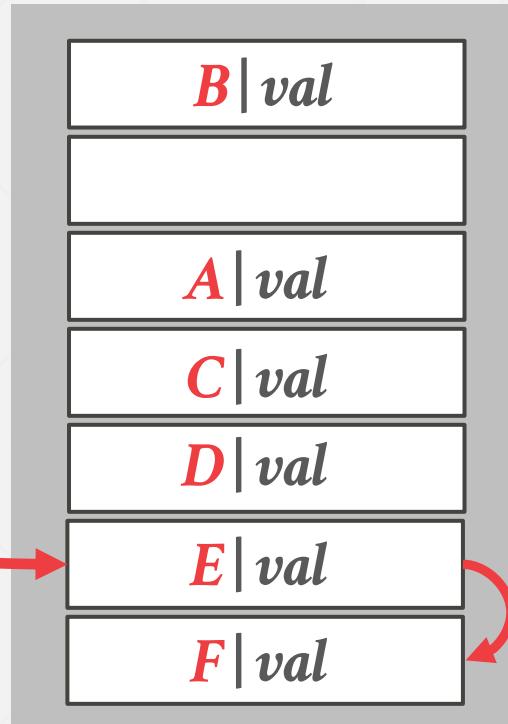
We go to

offset calc'

by hash fn
and starting

A
B
C
D
E

F



from there .

either until we find
Empty slot or

find the element.

So if delete removes

element causing

* white space then there

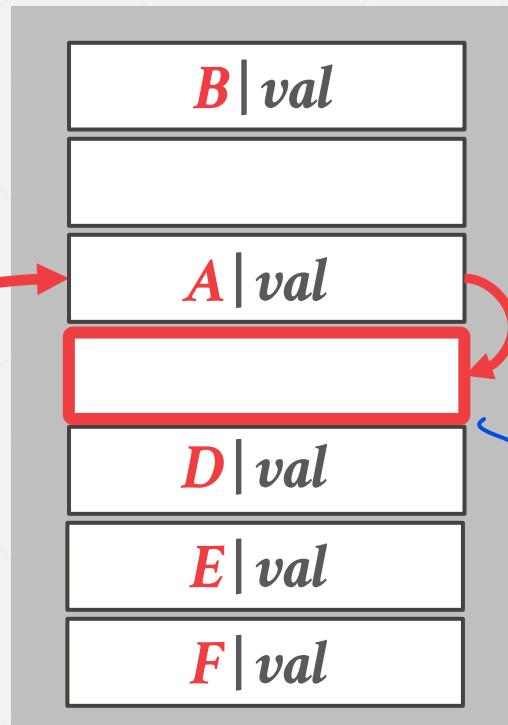
would be problem for

lookup ?

LINEAR PROBE HASHING - DELETES

$\text{hash}(\text{key}) \% N$

A
B
Delete C →
D
E
F

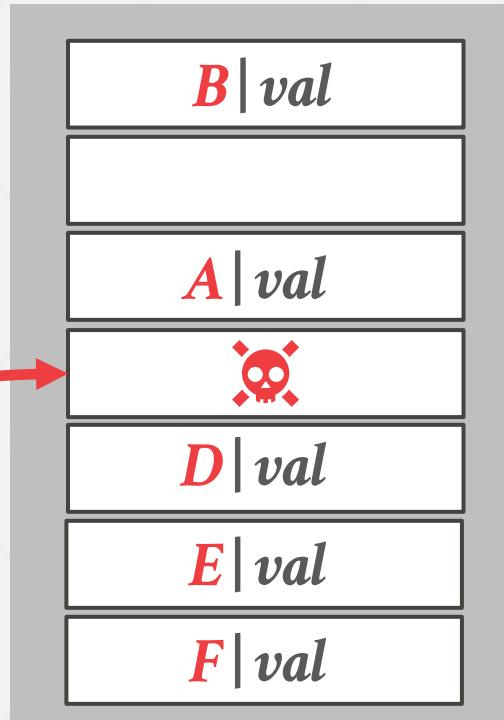


What happens during Delete?
we deleted c now we are searching for d.
as per hash it points here given in above slide.
But there we have empty slot - we stop search.

LINEAR PROBE HASHING - DELETES

$\text{hash}(\text{key}) \% N$

A
B
C
Get \rightarrow D
E
F

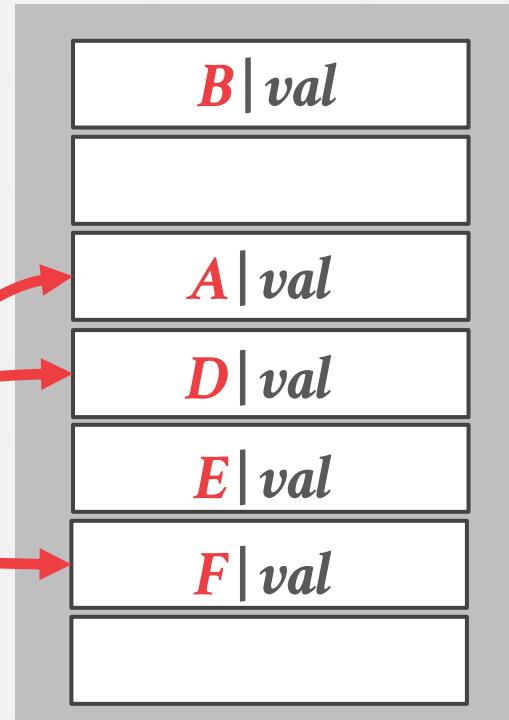
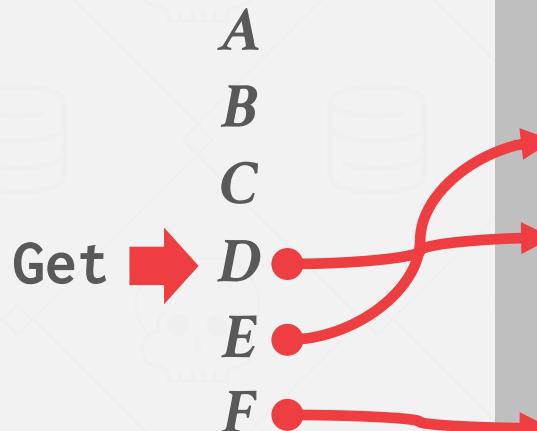


\therefore saying D not found but actually D is below the empty space.

\therefore Search gives wrong answer

LINEAR PROBE HASHING - DELETES

$\text{hash}(key) \% N$



Approach #1: Movement

- Rehash keys until you find the first empty slot.
- Nobody actually does this.

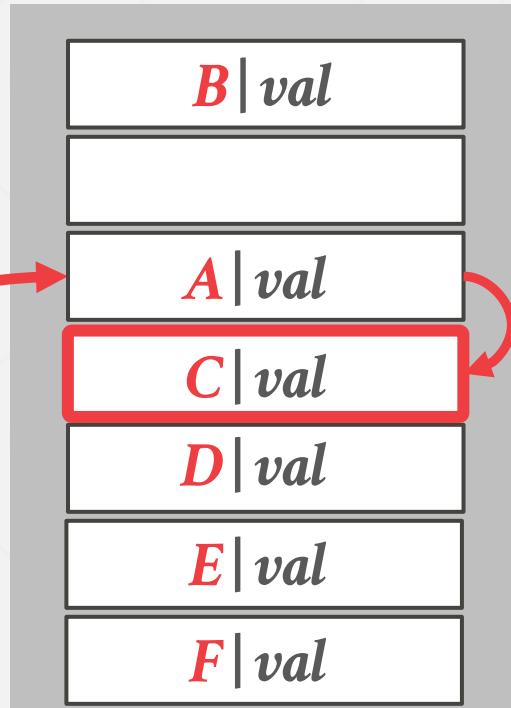
Move all keys D, E, F up. This is never preferred.

In tombstone we use marker to denote deleted keys. Later we can reuse them.

LINEAR PROBE HASHING - DELETES

$\text{hash}(\text{key}) \% N$

A
B
Delete \rightarrow C
D
E
F



deletes ↑ making many markers then we may have to

Approach #2: Tombstone

- Set a marker to indicate that the entry in the slot is logically deleted.
- You can reuse the slot for new keys.
- May need periodic garbage collection.

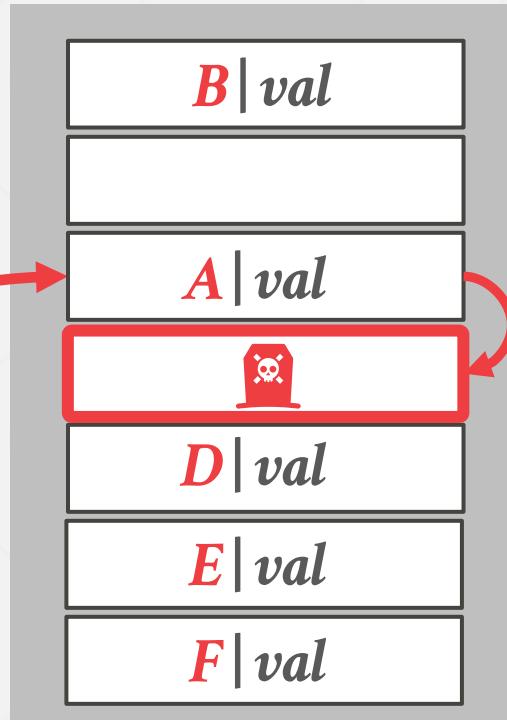
But returning is if re-off

do Garbage collection to collect them.

LINEAR PROBE HASHING - DELETES

$\text{hash}(key) \% N$

A
B
Delete 
C
D
E
F



Approach #2: Tombstone

- Set a marker to indicate that the entry in the slot is logically deleted.
- You can reuse the slot for new keys.
- May need periodic garbage collection.

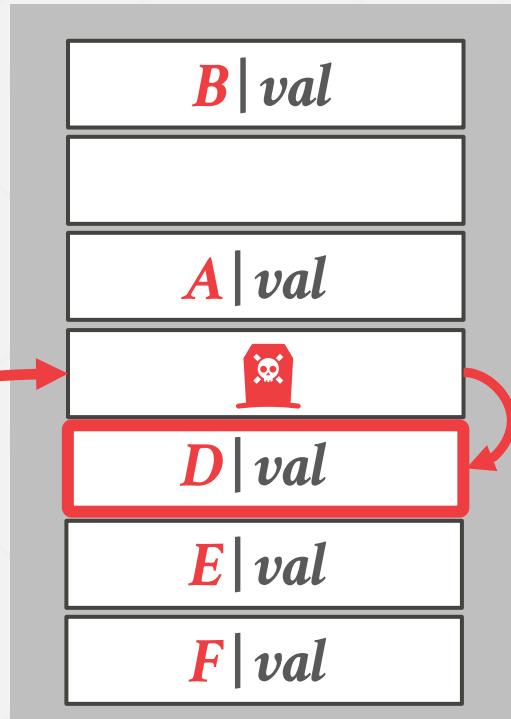
In this case at search
in certain slot if we

find marker then we just move next to it we don't stop there.

LINEAR PROBE HASHING - DELETES

$\text{hash}(\text{key}) \% N$

A
B
C
Get \rightarrow D
E
F



Approach #2: Tombstone

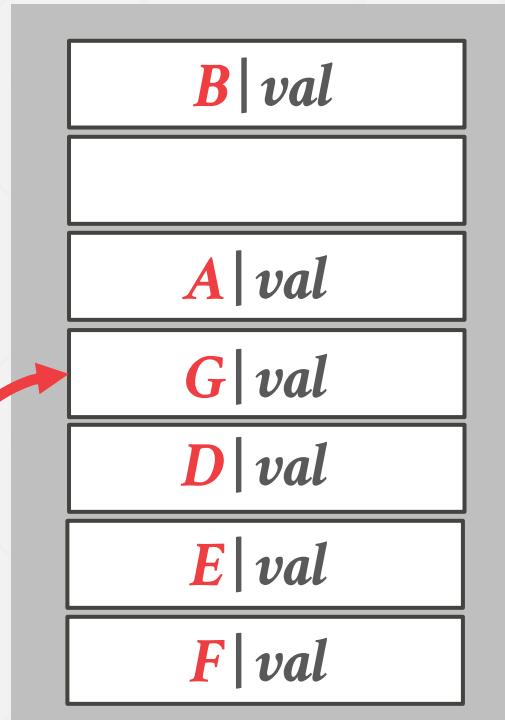
- Set a marker to indicate that the entry in the slot is logically deleted.
- You can reuse the slot for new keys.
- May need periodic garbage collection.

LINEAR PROBE HASHING - DELETES

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F
G

Put



Approach #2: Tombstone

- Set a marker to indicate that the entry in the slot is logically deleted.
- You can reuse the slot for new keys.
- May need periodic garbage collection.

If a key has more than 1 value

then we store all

NON-UNIQUE KEYS

those values of single key in a separate linked list.

Choice #1: Separate Linked List

- Store values in separate storage area for each key.

What are the examples?

Choice #2: Redundant Keys

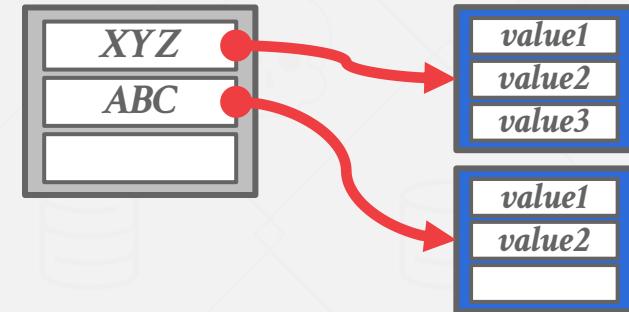
- Store duplicate keys entries together in the hash table.
- This is easier to implement so this is what most systems do.

Redundant ⇒ Duplicate keys. Store duplicate

keys

together in hash table

Value Lists



XYZ value2
ABC value1
XYZ value3
XYZ value1
ABC value2

A variant of linear probe hashing that steals slots from rich keys &

gives them to

ROBIN HOOD HASHING

poor keys.

Variant of linear probe hashing that steals slots from "rich" keys and give them to "poor" keys.

- Each key tracks the number of positions they are from where its optimal position in the table. → i.e. we find expected position from hash fn. Let it be a .
- On insert, a key takes the slot of another key if the first key is farther away from its optimal position than the second key.

But it is some where at b .

$(a-b)$ is larger than

the 2nd key's $(a-b)$.

∴ we keep track of this

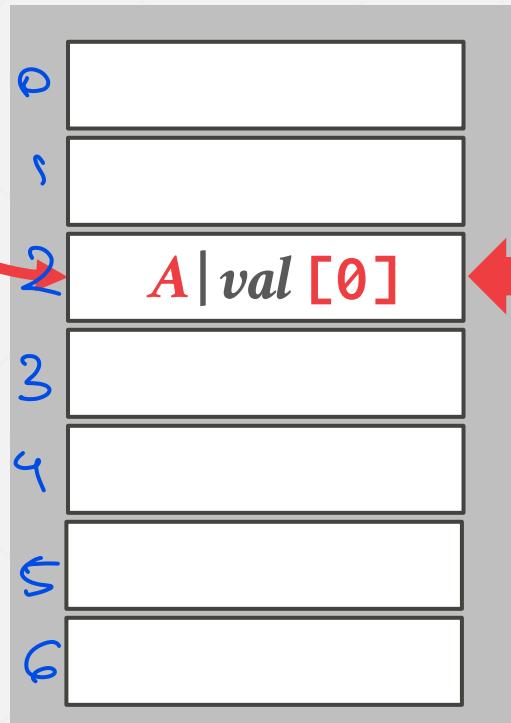
$(a-b)$ for

all keys.

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



A q's expected to stay at 2

∴ Empty place $M \vdash$

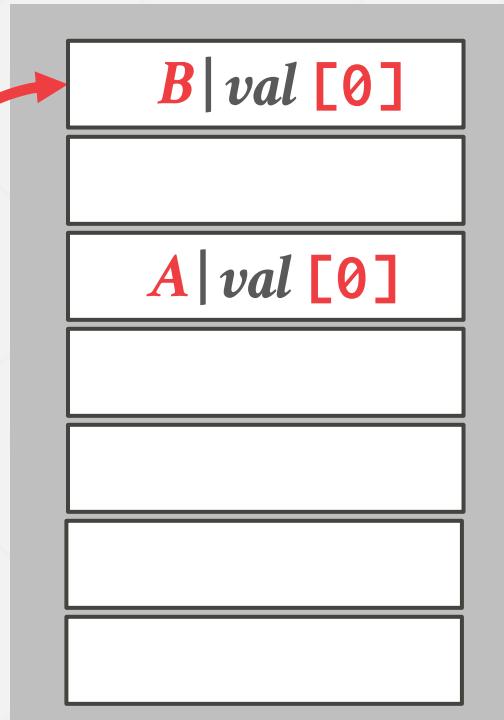
$$\Lambda(a-b) = 0$$

of "Jumps" From First Position

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



Place B → as per fn,

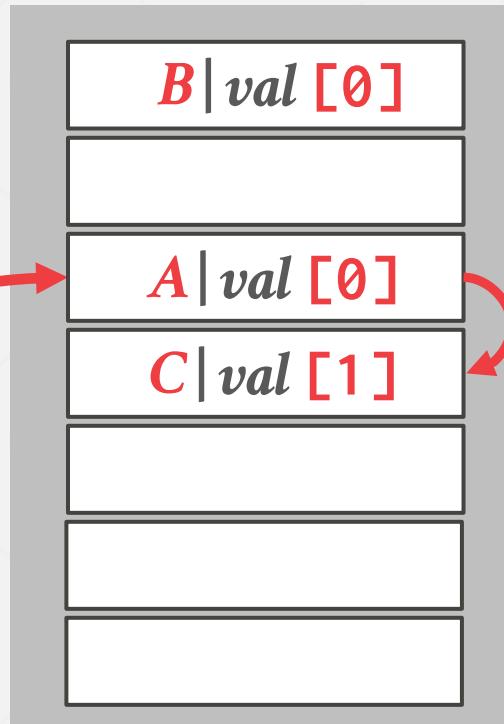
B should be at 0 it is

∴ B can't = 0

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



C should be at 2 but not there

$A[0] == C[0]$
 $\therefore \text{Insert}(C)$ is happening.

C takes position of A

if $C(a-b) > A(a-b)$

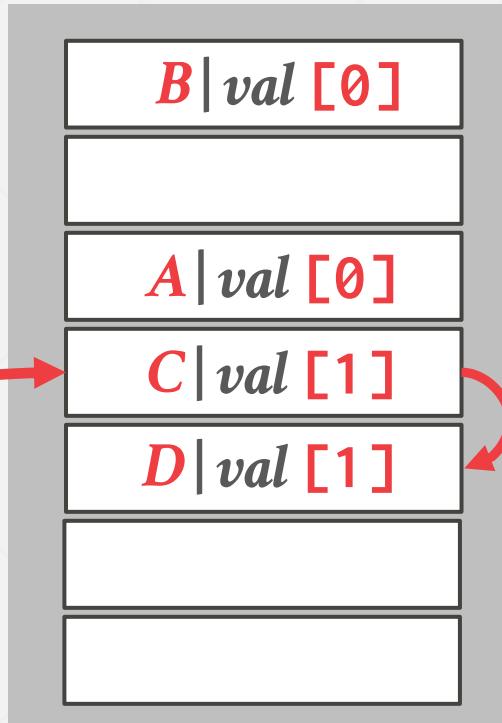
both zero. \therefore no

$C(a-b)++$ C goes down

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



My increment D's (a-b)

Insert C Now

Insert D.

$D(a-b) = 0$ -

Collision key \rightarrow C

$C[1] > D[0]$ $C(a-b) = 1$

$0 > 1$

... D goes down \because it

Cannot force C's position

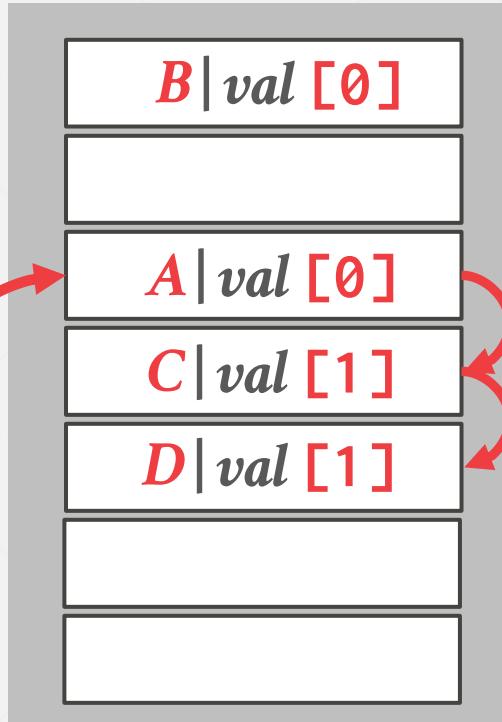
ROBIN HOOD HASHING

Insert E.

Collision $\rightarrow A$

$hash(key) \% N$

A
B
C
D
E
F



Subsequently $D(a-b)$ also changes

$E(a-b) \rightarrow 1$
Collision $\rightarrow C$

$A[0] == E[0]$

$C[1] == E[1]$

$D[1] < E[2]$ Collision $\rightarrow D$.

$2 > 1 \therefore E$ takes

position of D

D moves down

$E(a-b) \geq 0$ $A(a-b) = 0$
 $0 > 0 \times$

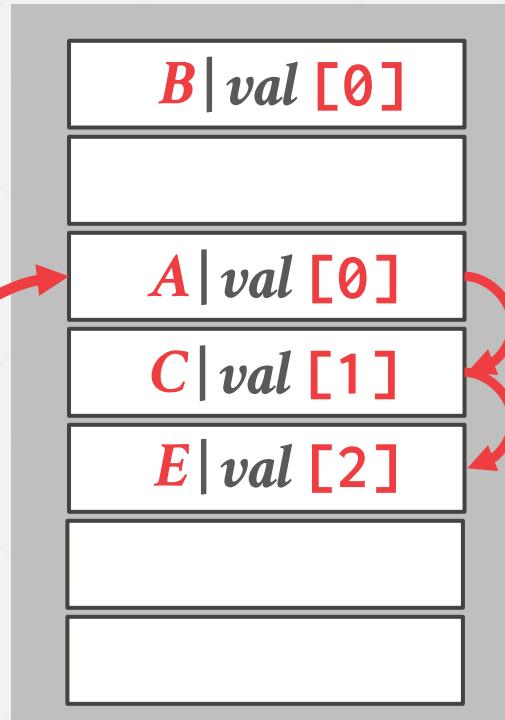
$I \geq 1 \wedge$

$E(a-b) \rightarrow 2$

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



$A[0] == E[0]$

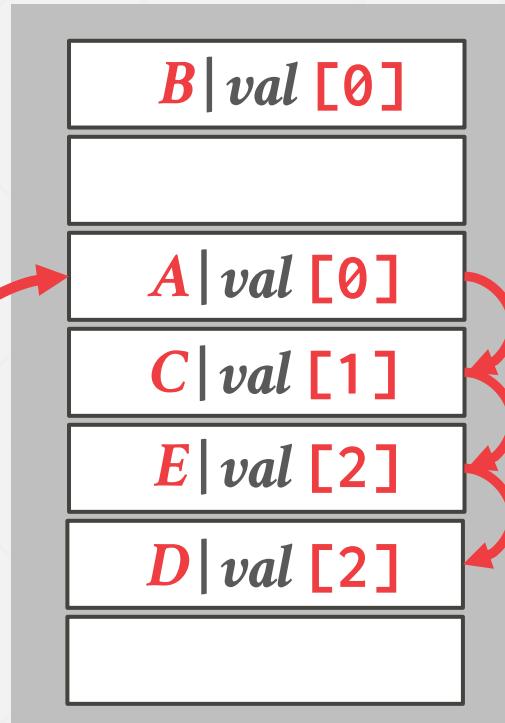
$C[1] == E[1]$

$D[1] < E[2]$

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A
B
C
D
E
F



$A[0] == E[0]$

$C[1] == E[1]$

$D[1] < E[2]$

ROBIN HOOD HASHING

$\text{hash}(\text{key}) \% N$

A

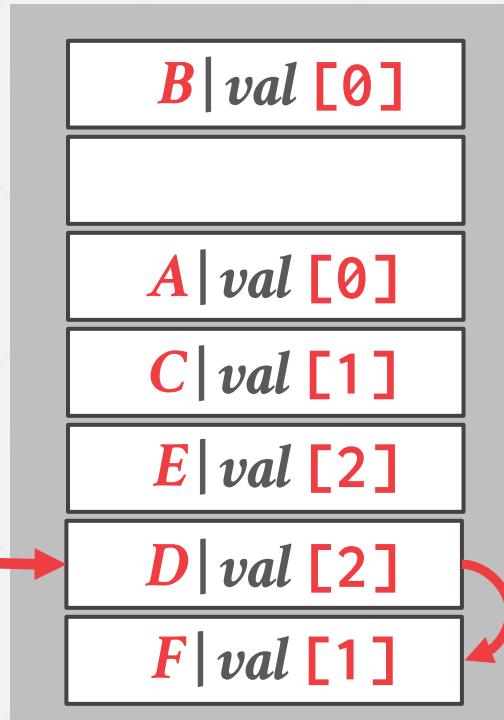
B

C

D

E

F



We use multiple hash tables ↗ Each table has its own hash function.

CUCKOO HASHING

On insert check

every table &

pick up the
available
free slot.

Use multiple hash tables with different hash function seeds.

- On insert, check every table and pick anyone that has a free slot.
- If no table has a free slot, evict the element from one of them and then re-hash it find a new location.

Look-ups and deletions are always **$O(1)$** because only one location per hash table is checked.

Best open-source implementation is from CMU.

↳ Only one location in Each Table is checked.

If no table

has free slot,

evict an
element from
one of them

& place it
there.

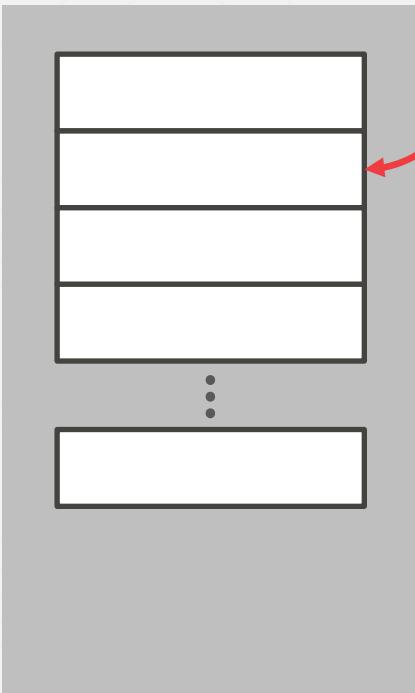
Removed Ele

needs to be
rehashed to
new location



CUCKOO HASHING

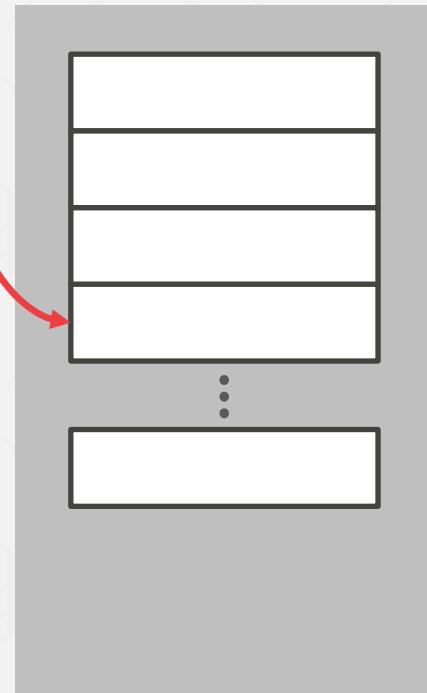
Hash Table #1



Put A

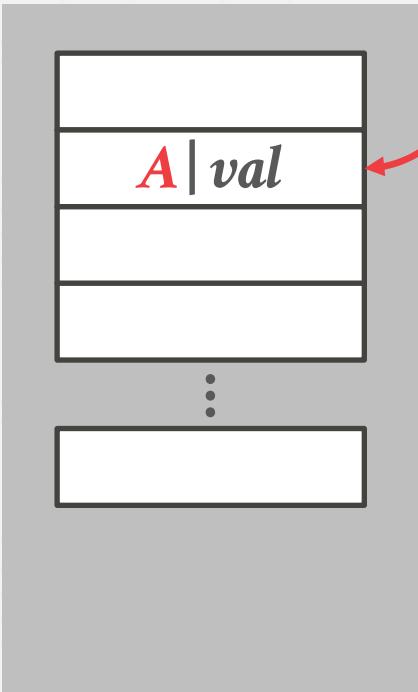
$hash_1(A)$ $hash_2(A)$

Hash Table #2



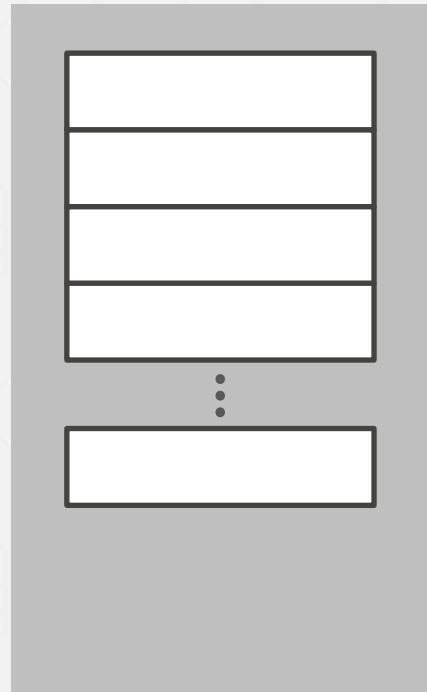
CUCKOO HASHING

Hash Table #1



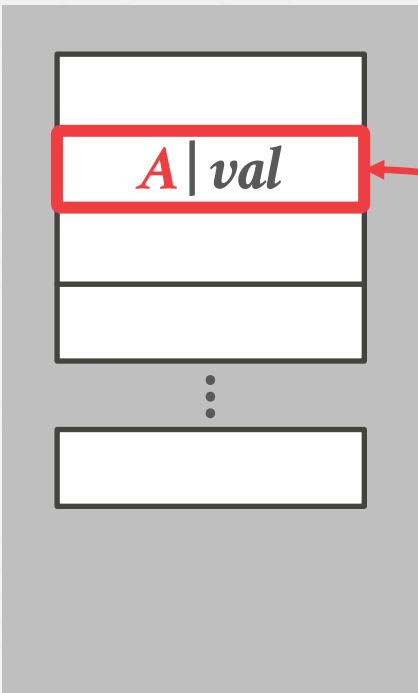
Put A
 $hash_1(A)$ $hash_2(A)$
:- we placed A at free slot
inTable1 -

Hash Table #2



CUCKOO HASHING

Hash Table #1



Put A
 $hash_1(A)$ $hash_2(A)$

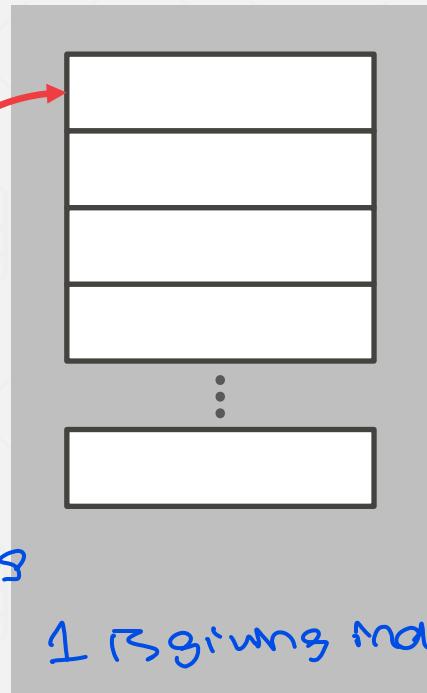
Put B
 $hash_1(B)$ $hash_2(B)$

Now put B free slot not M

Table1 :: put in Table2

Both hashes are not giving
some offset

Hash Table #2

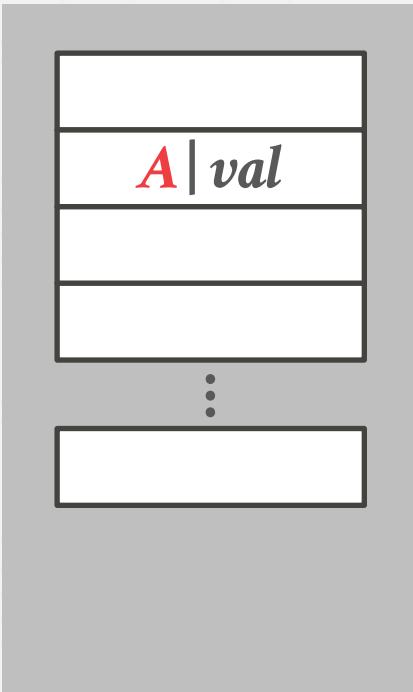


1 is giving index2

{ 2 is giving index1

CUCKOO HASHING

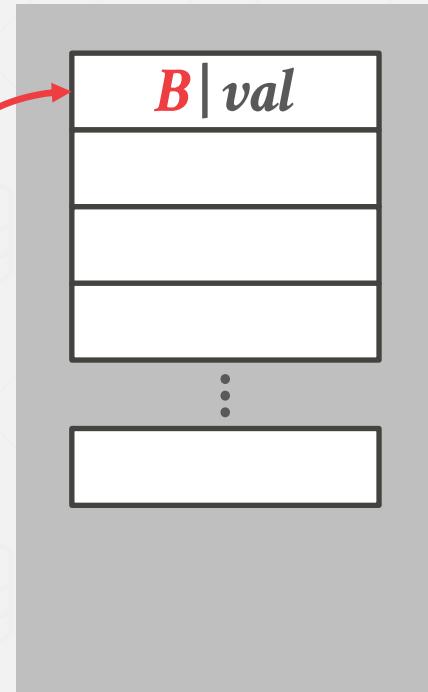
Hash Table #1



Put A
 $hash_1(A)$ $hash_2(A)$

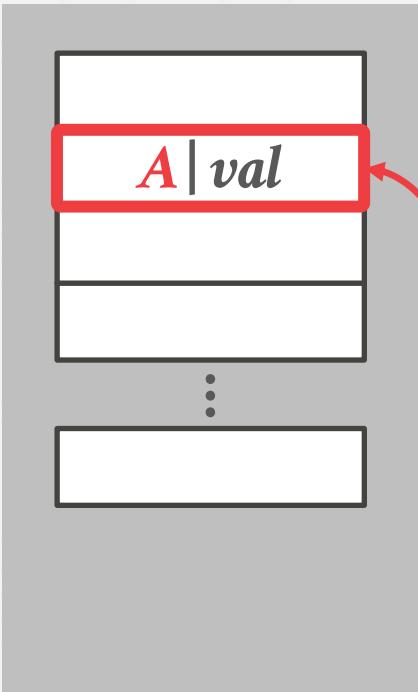
Put B
 $hash_1(B)$ $hash_2(B)$

Hash Table #2



CUCKOO HASHING

Hash Table #1

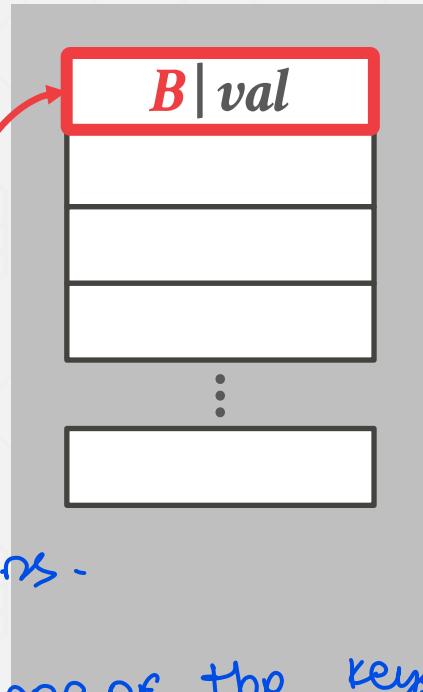


Put A
 $hash_1(A)$ $hash_2(A)$

Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$

Hash Table #2



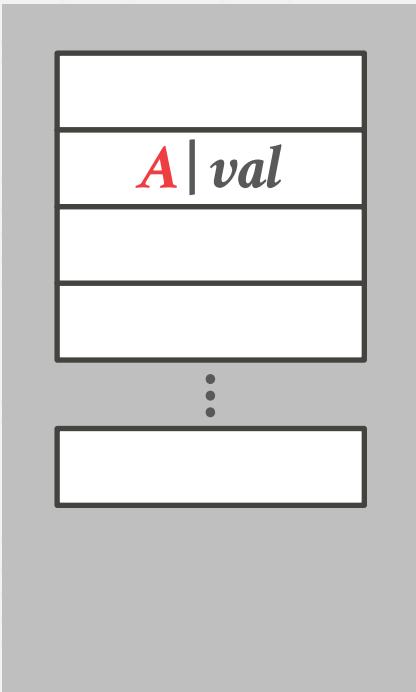
Now for C both are

mapping to filled locations -

-- we need to evict any one of the keys -

CUCKOO HASHING

Hash Table #1

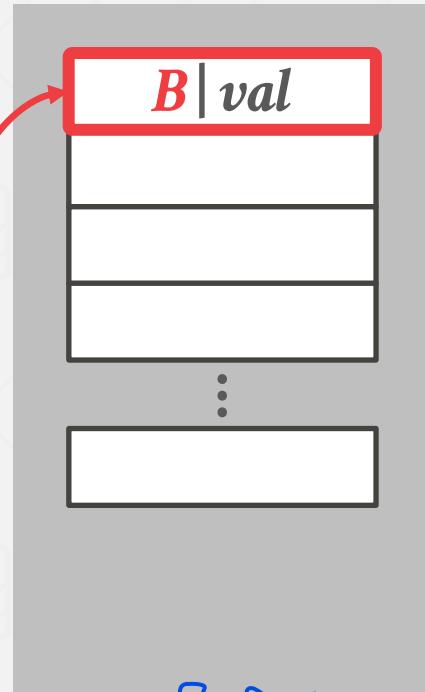


Put A
 $hash_1(A)$ $hash_2(A)$

Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$

Hash Table #2

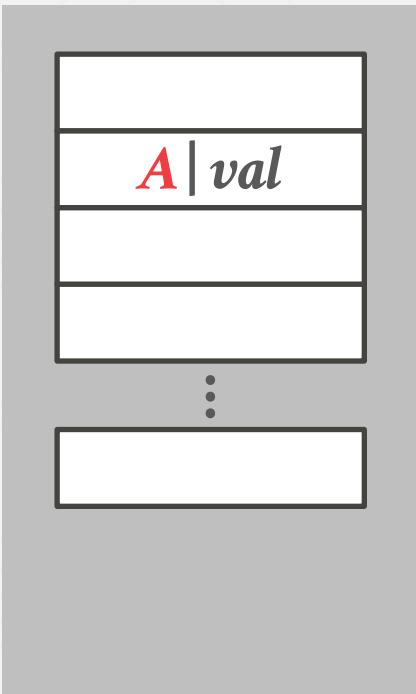


We need to evict either
 A or B to place C.

Then only O(1) for Search & Delete.

CUCKOO HASHING

Hash Table #1

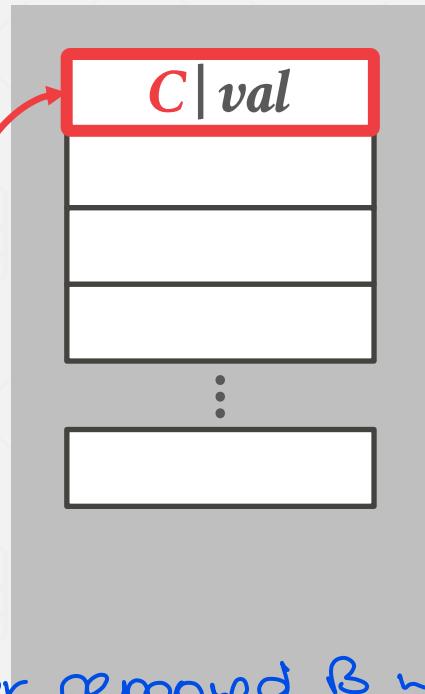


Put A
 $hash_1(A)$ $hash_2(A)$

Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$

Hash Table #2



∴ we removed B and
 placed C.

O(1) and for removed B we
 need to rehash and

$\therefore B$ is removed
from table2

place it in new
location

CUCKOO HASHING

Hash Table #1

C	val
A	val
⋮	

Put A
 $hash_1(A)$ $hash_2(A)$

Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$
 $hash_1(B)$

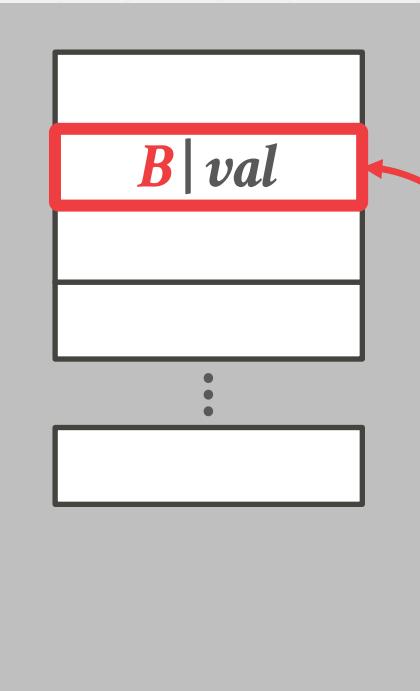
Hash Table #2

C	val
C	
⋮	

place B at index obtained
from hashing at
table1.

CUCKOO HASHING

Hash Table #1

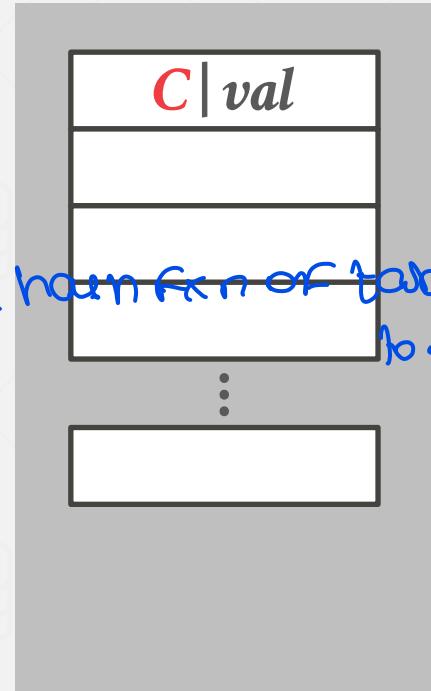


Put A
 $hash_1(A)$ $hash_2(A)$

Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$
 $hash_1(B)$

Hash Table #2



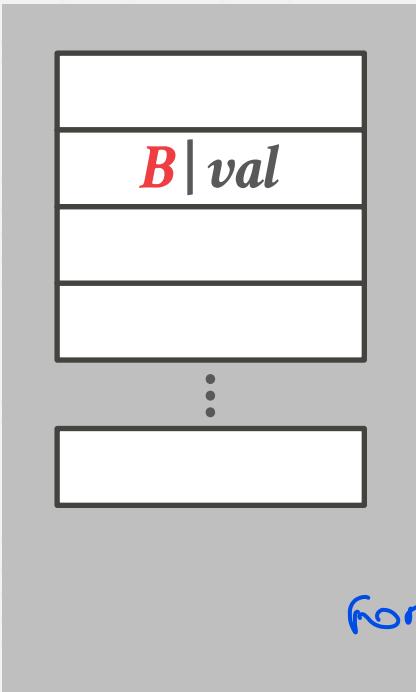
use hash fn of table 2
to place A

But B is now colliding with A.

- = Evict A
place B
rehash

CUCKOO HASHING

Hash Table #1



Put A
 $hash_1(A)$ $hash_2(A)$

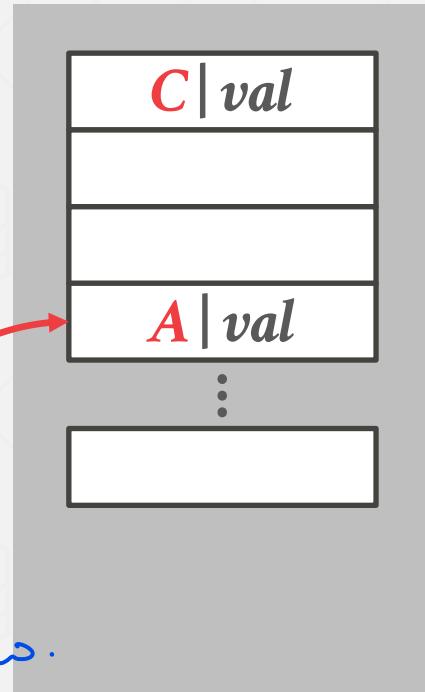
Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$

$hash_1(B)$ ←
 $hash_2(A)$ →
 Seq OF steps

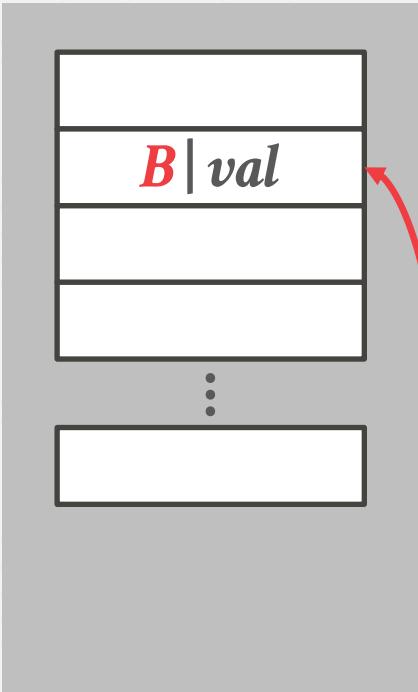
∴ we placed A in
 table 2 now.

Hash Table #2



CUCKOO HASHING

Hash Table #1



Put A
 $hash_1(A)$ $hash_2(A)$

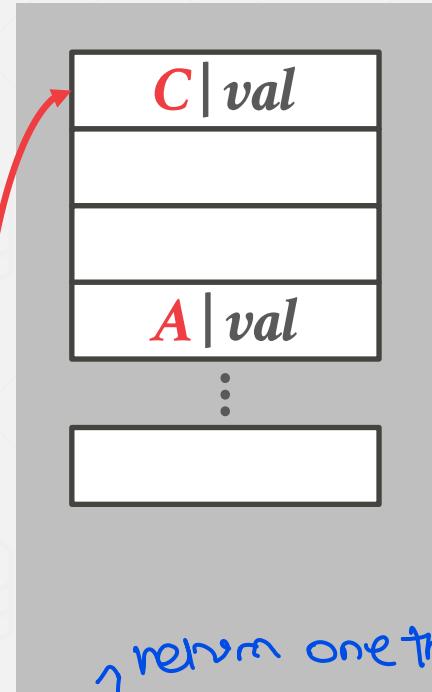
Put B
 $hash_1(B)$ $hash_2(B)$

Put C
 $hash_1(C)$ $hash_2(C)$
 $hash_1(B)$
 $hash_2(A)$

Get B
 $hash_1(B)$ $hash_2(B)$

Now $\text{Search}(B) \rightarrow 2$ offsets

Hash Table #2



↑ return one that
has 0

OBSERVATION

The previous hash tables require the DBMS to know the number of elements it wants to store.

- Otherwise, it must rebuild the table if it needs to grow/shrink in size.

All these req.no.of elements we have to store.

We simply cannot maintain more than 1 hash table

Dynamic hash tables resize themselves on demand.

- Chained Hashing
- Extendible Hashing
- Linear Hashing

or a giant table.

Each slot in hash table
is a linked list
of buckets.

CHAINED HASHING

For Search we
hash to bucket

& scan it

Maintain a linked list of buckets for each slot in the hash table.

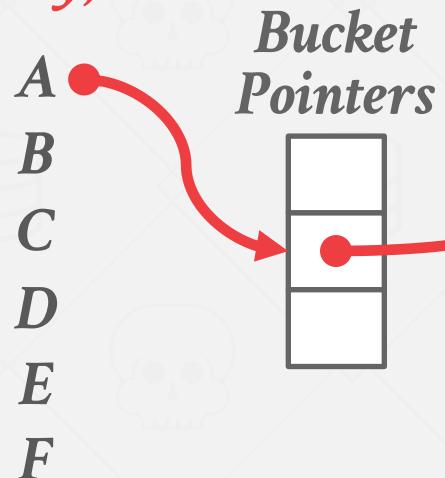
Resolve collisions by placing all elements with the same hash key into the same bucket.

- To determine whether an element is present, hash to its bucket and scan for it.
- Insertions and deletions are generalizations of lookups.

all elements with some hash key into some bucket.

CHAINED HASHING

$\text{hash}(\text{key}) \% N$

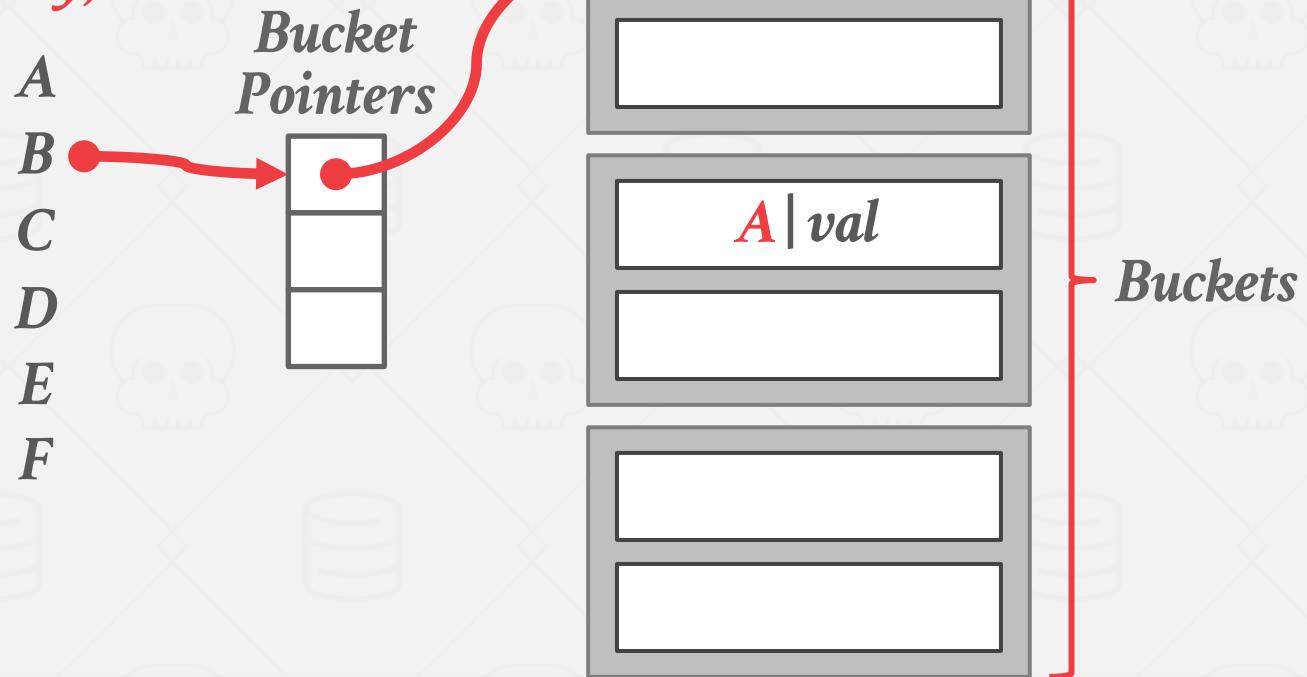


To maintain buckets we have list to store references to Buckets.

↳ a linked list.
Buckets

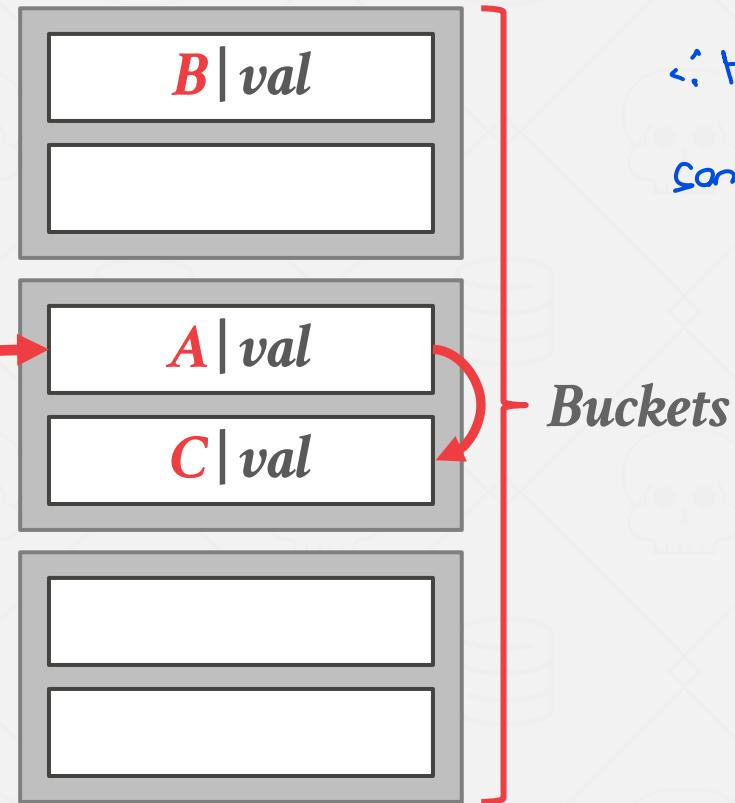
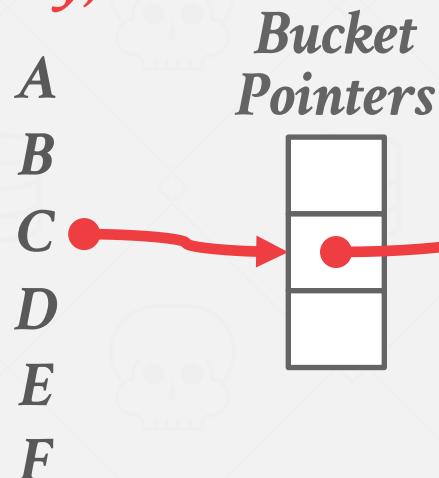
CHAINED HASHING

$hash(key) \% N$



CHAINED HASHING

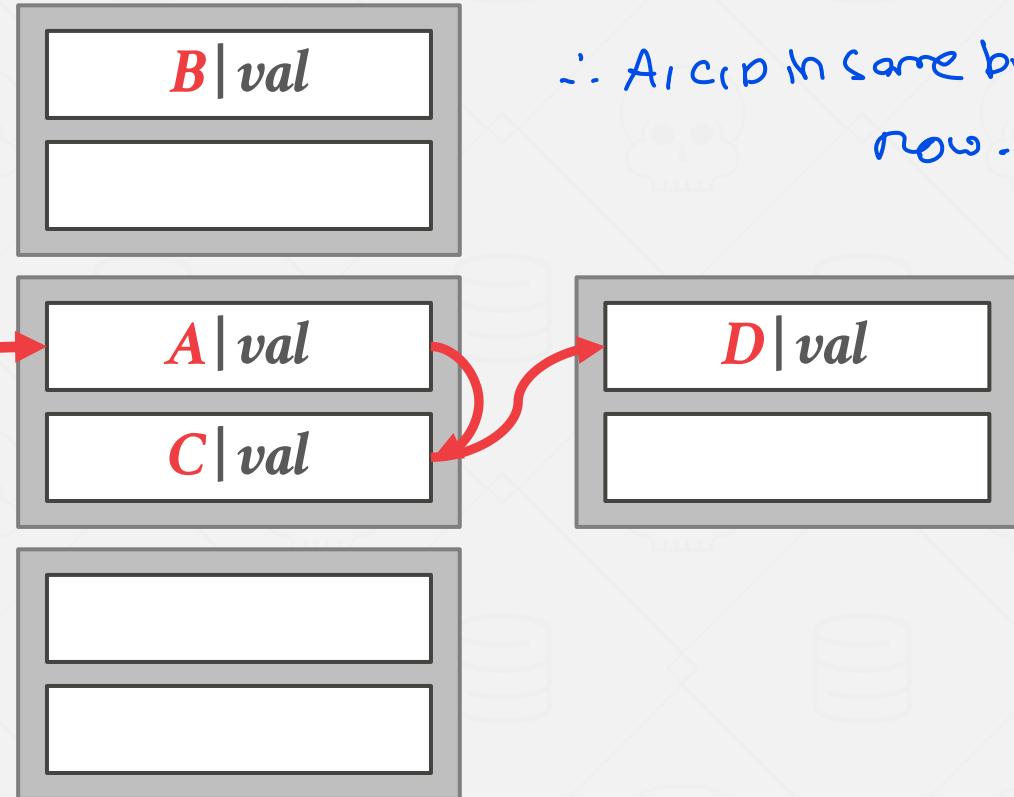
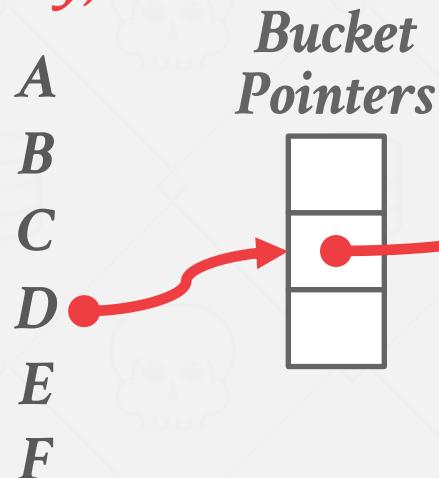
$\text{hash}(\text{key}) \% N$



$\because \text{Here } A \in C$
 same key \therefore They
 both set into
 some bucket.

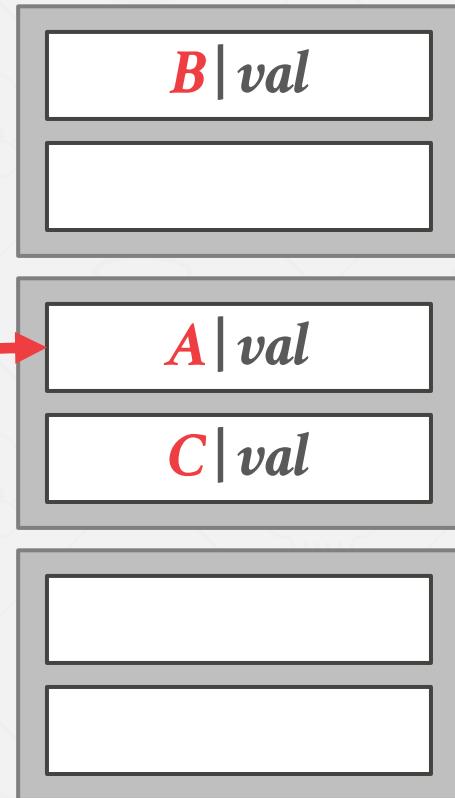
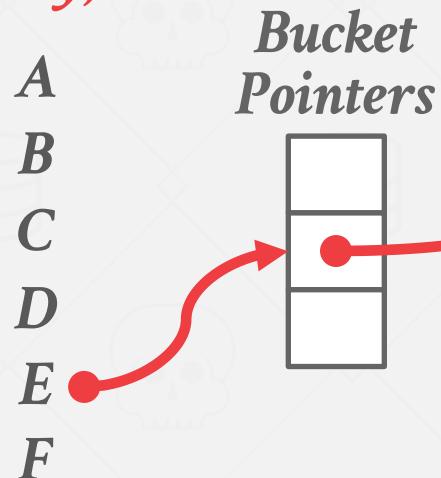
CHAINED HASHING

$\text{hash}(\text{key}) \% N$



CHAINED HASHING

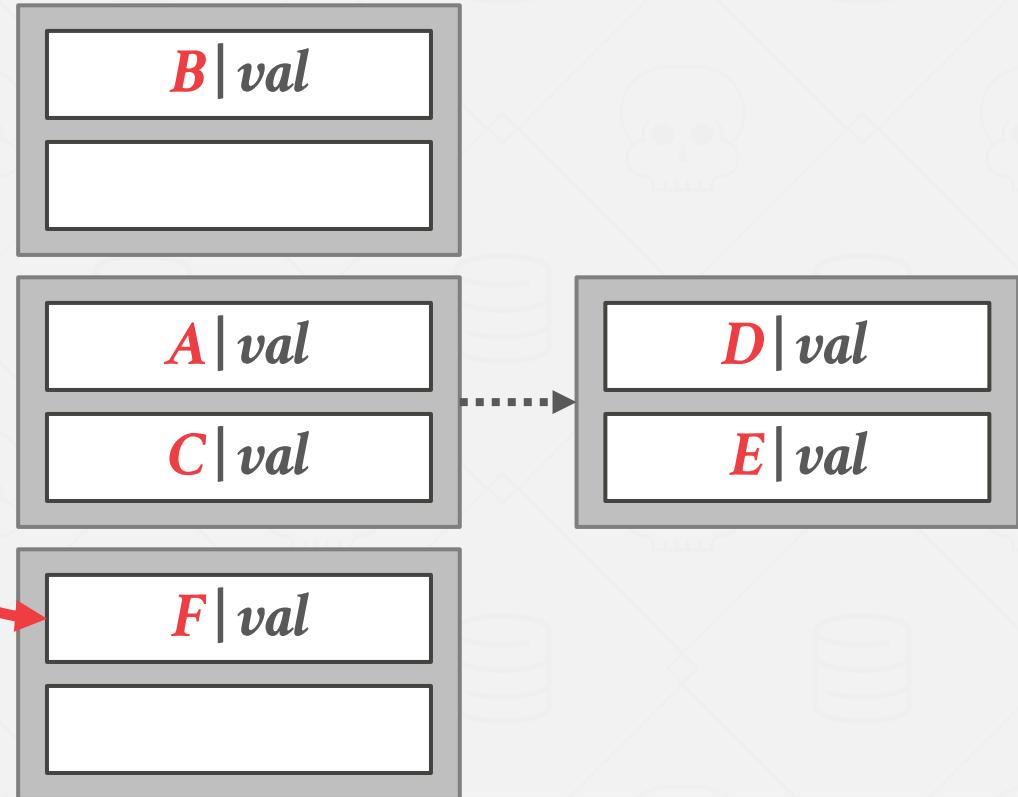
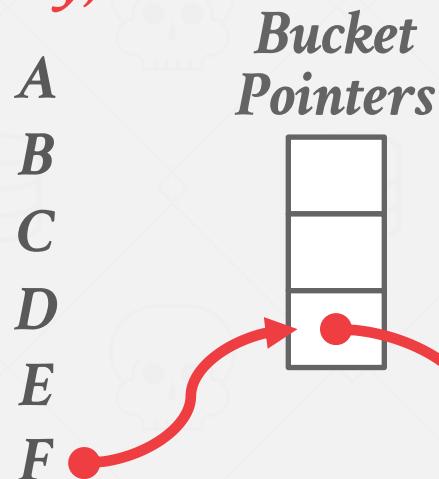
$\text{hash}(\text{key}) \% N$



Also hashed to same
point ∵ A, C, D, E all
stored are ~~pointed~~ to
some bucket

CHAINED HASHING

$hash(key) \% N$



Variant of

Chained hashing

Where we split

buckets instead of letting linked lists grow forever.

EXTENDIBLE HASHING

Chained-hashing approach where we split buckets instead of letting the linked list grow forever.

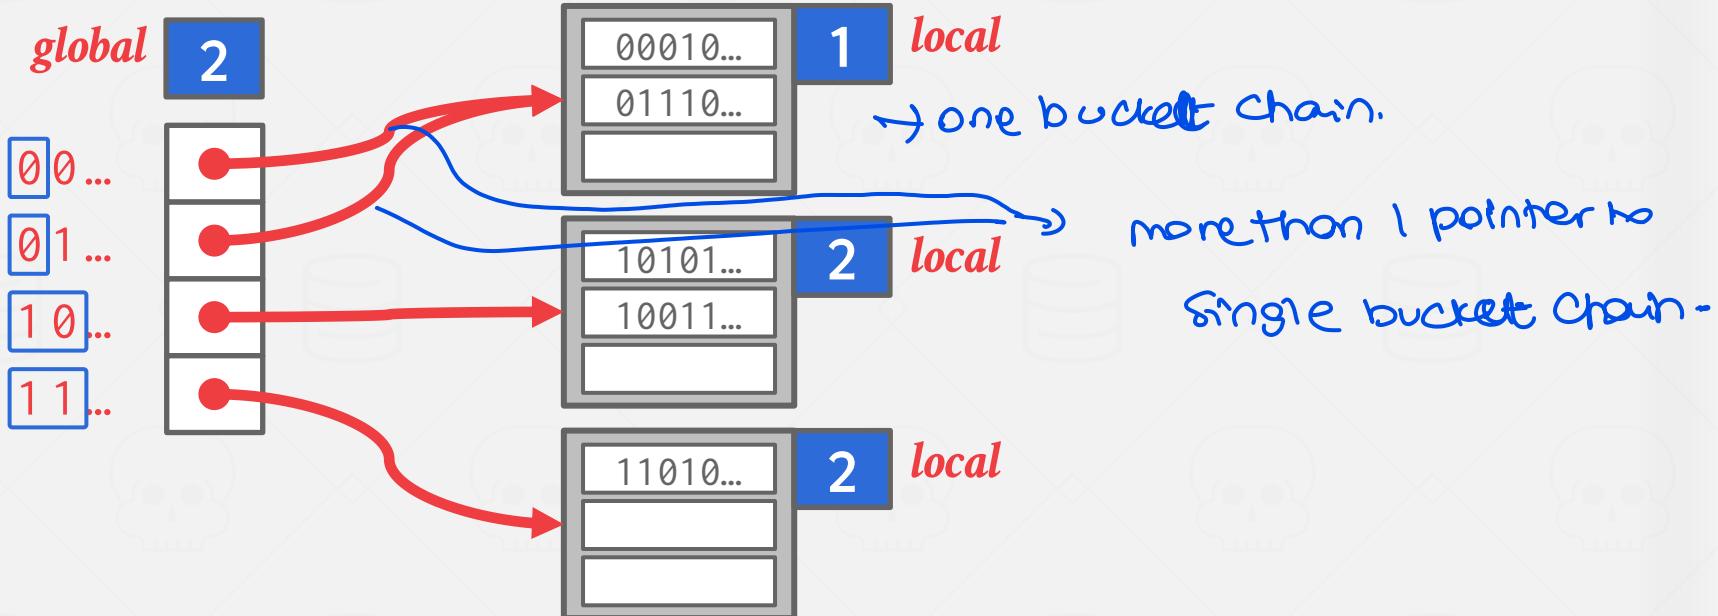
Multiple slot locations can point to the same
bucket chain.

Reshuffle bucket entries on split and increase the number of bits to examine.

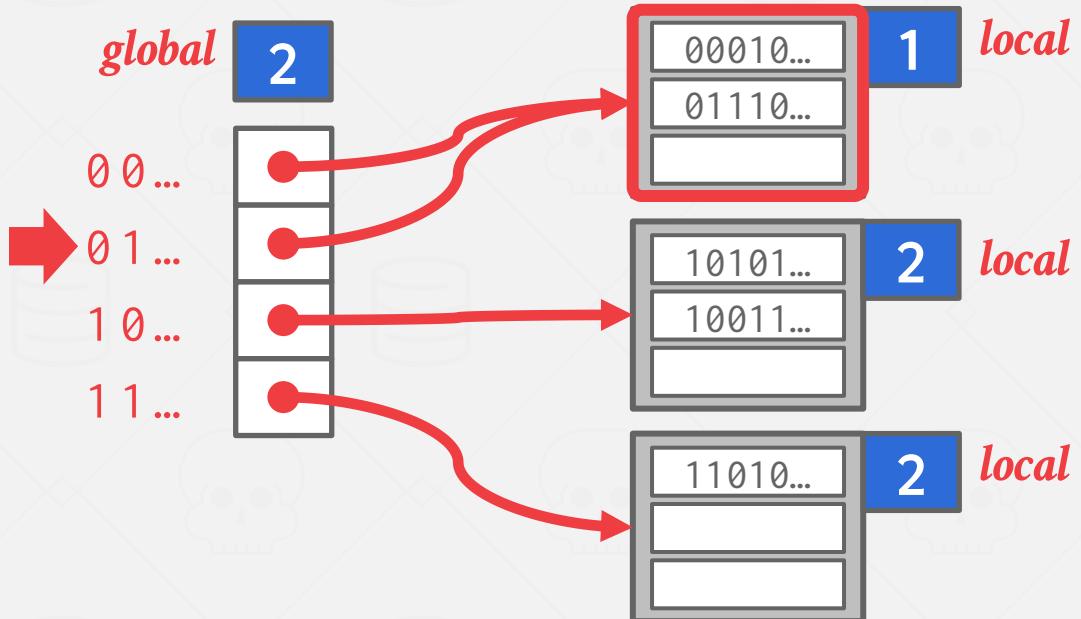
→ Data movement is localized to just the split chain.

During split reshuffle the bucket entries & increase the no. of bits examined

EXTENDIBLE HASHING



EXTENDIBLE HASHING



Get A
 $\text{hash}(A) = 01110\ldots$

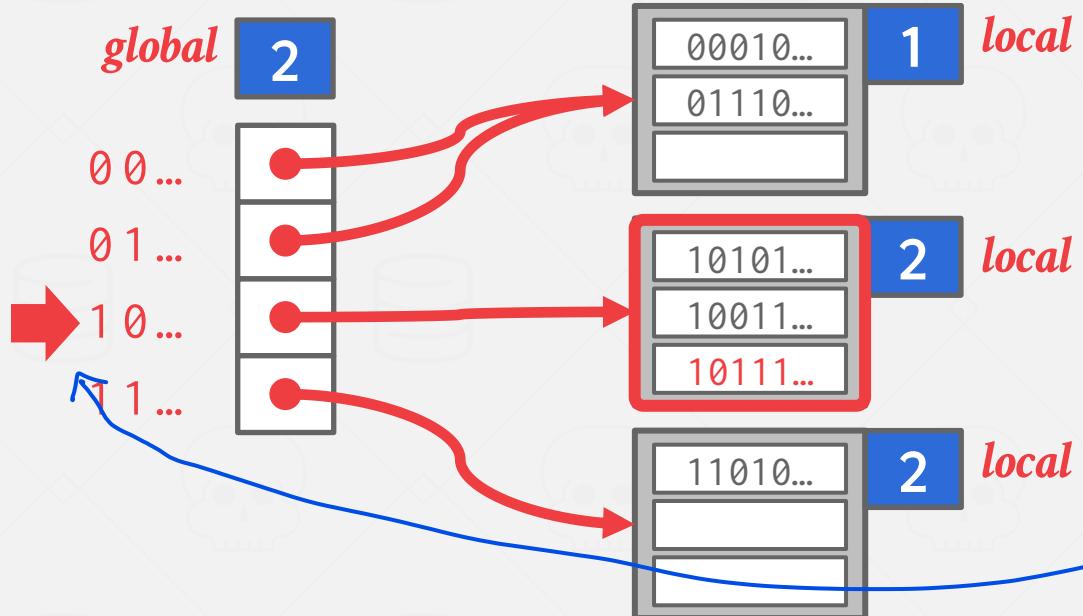
Upon Lookup.

∴ Go to bucket chain 1

? scan for it

way to check
for offset

EXTENDIBLE HASHING



Get A
 $\text{hash}(A) = 01110...$

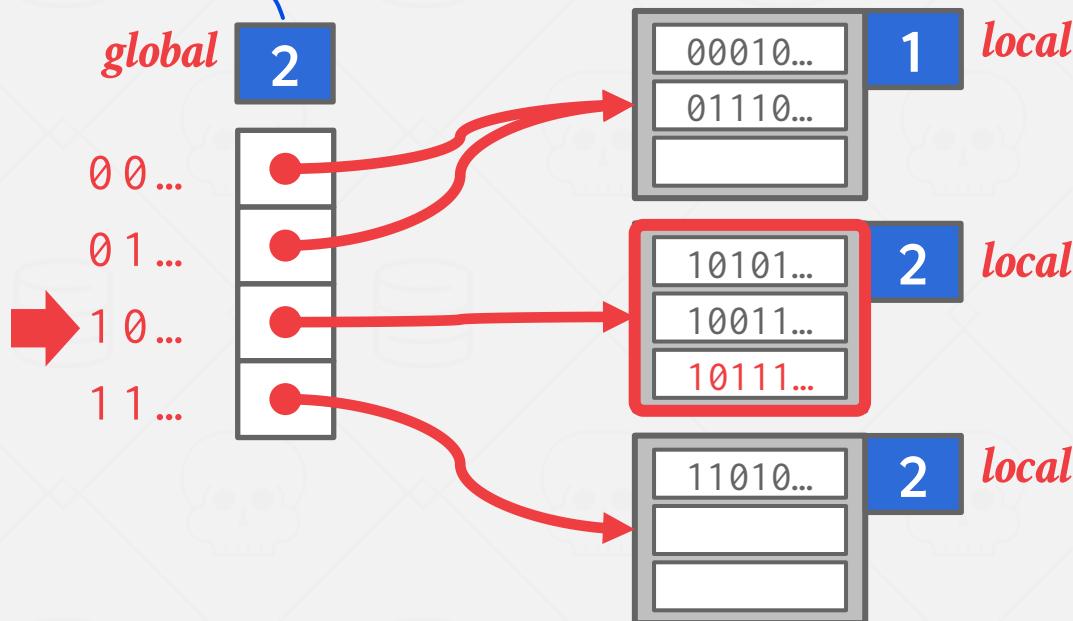
Put B
 $\text{hash}(B) = 10111...$

\therefore Go to bucket chain

Place B there.

2

2 bits are req.



Get A

$\text{hash}(A) = 01110\ldots$

Put B

$\text{hash}(B) = 10111\ldots$

Put C

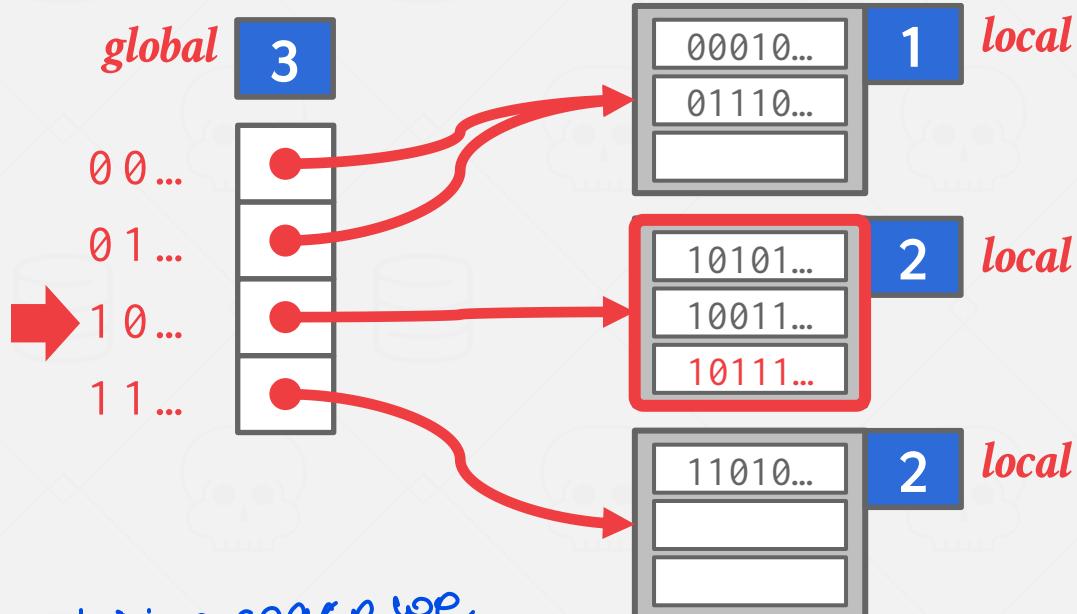
$\text{hash}(C) = \boxed{10100}\ldots$



go to bucket chain 2

check for empty space

EXTENDIBLE HASHING



$\therefore \because$ no space we
need to split & increase the
no. of bits for
lookup.

Get A
 $hash(A) = 01110...$

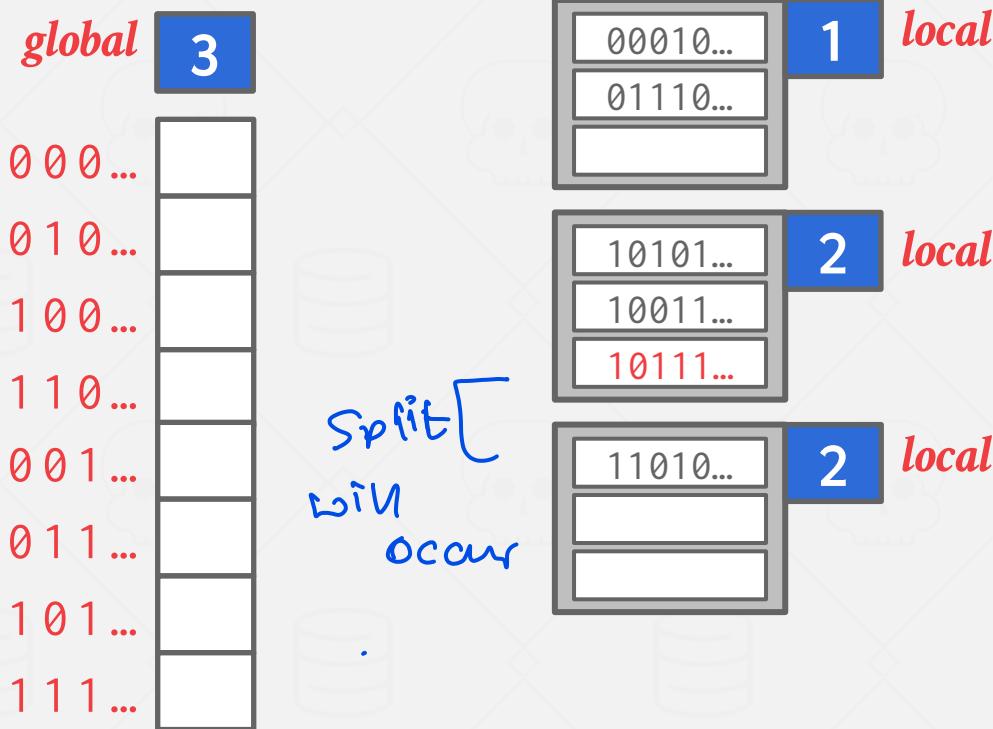
Put B
 $hash(B) = 10111...$

Put C
 $hash(C) = 10100...$

C does not have space to be
put in bucket chain 2

- $2 \rightarrow 3$ if split occurs if entries are reshuffled-

EXTENDIBLE HASHING

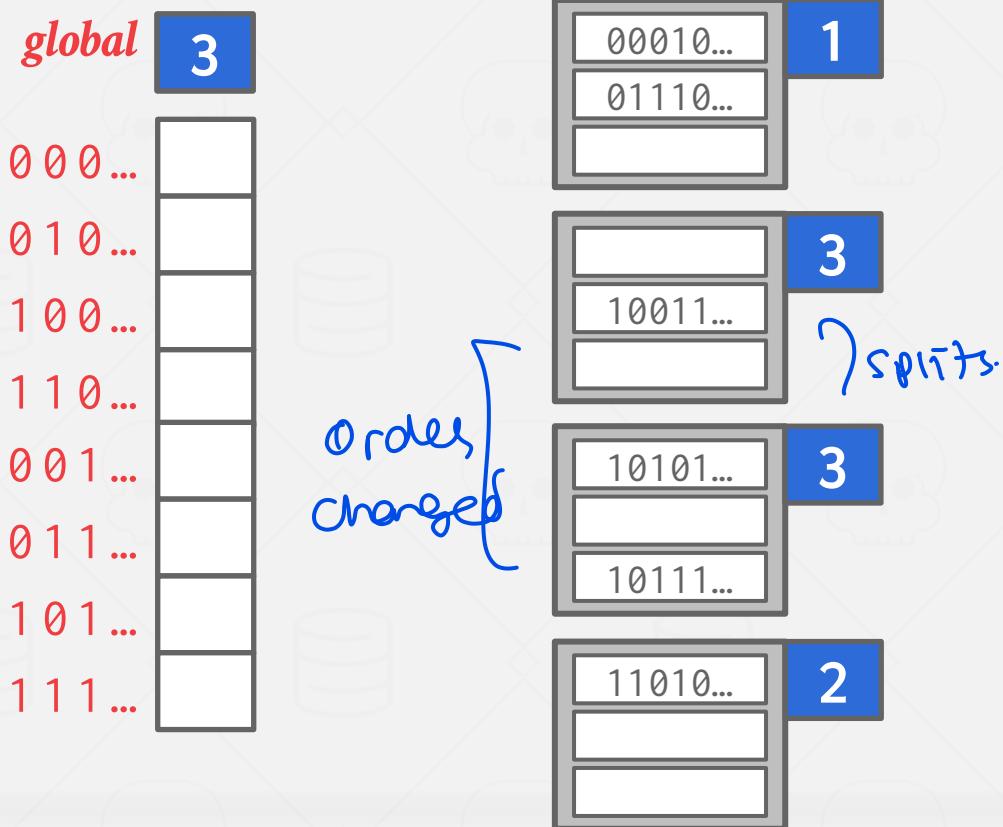


Get A
 $\text{hash}(A) = 01110...$

Put B
 $\text{hash}(B) = 10111...$

Put C
 $\text{hash}(C) = 10100...$

EXTENDIBLE HASHING



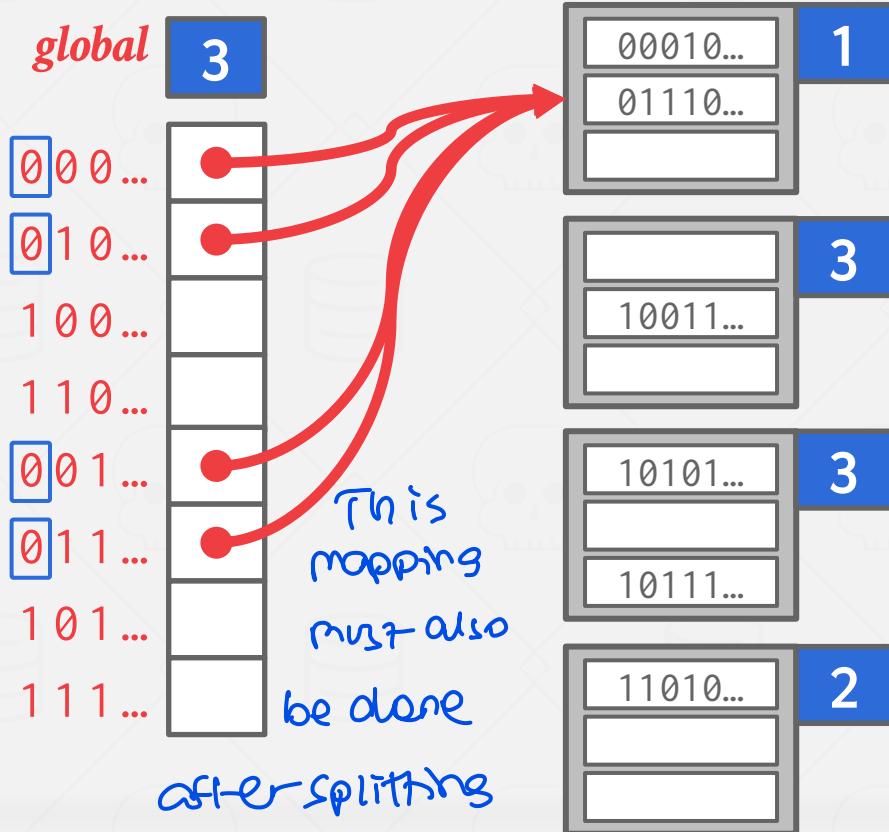
Get A
 $\text{hash}(A) = 01110...$

Put B
 $\text{hash}(B) = 10111...$

Put C
 $\text{hash}(C) = 10100...$

now we put C

EXTENDIBLE HASHING



Get A

$\text{hash}(A) = 01110\ldots$

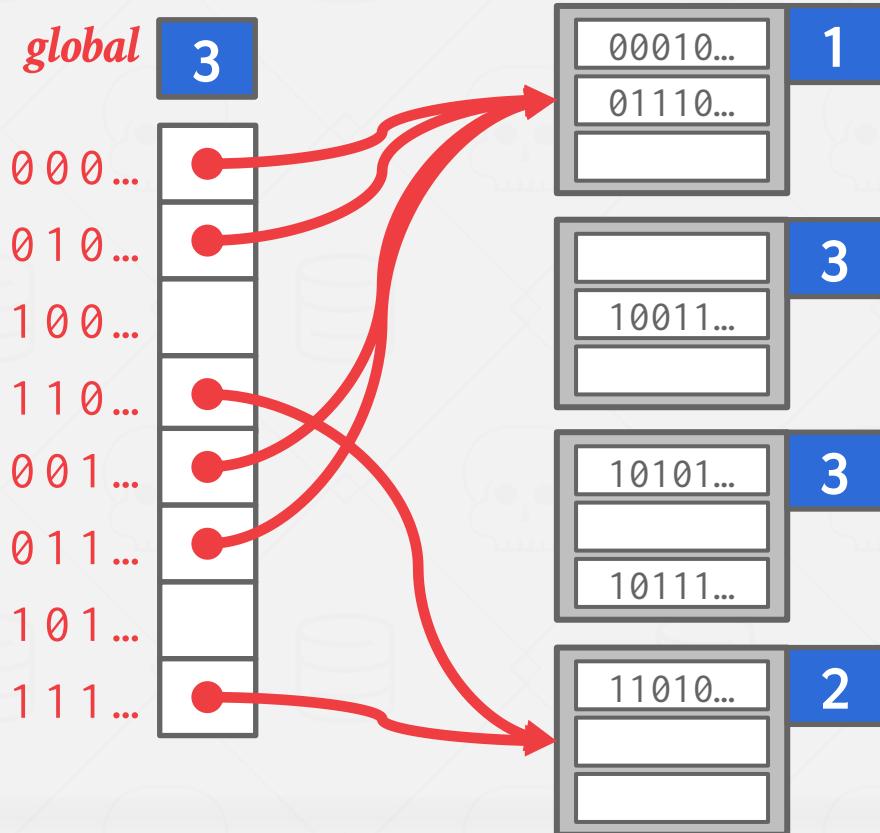
Put B

$\text{hash}(B) = 10111\ldots$

Put C

$\text{hash}(C) = \underline{10100}\ldots$

EXTENDIBLE HASHING



Get A

$\text{hash}(A) = 01110...$

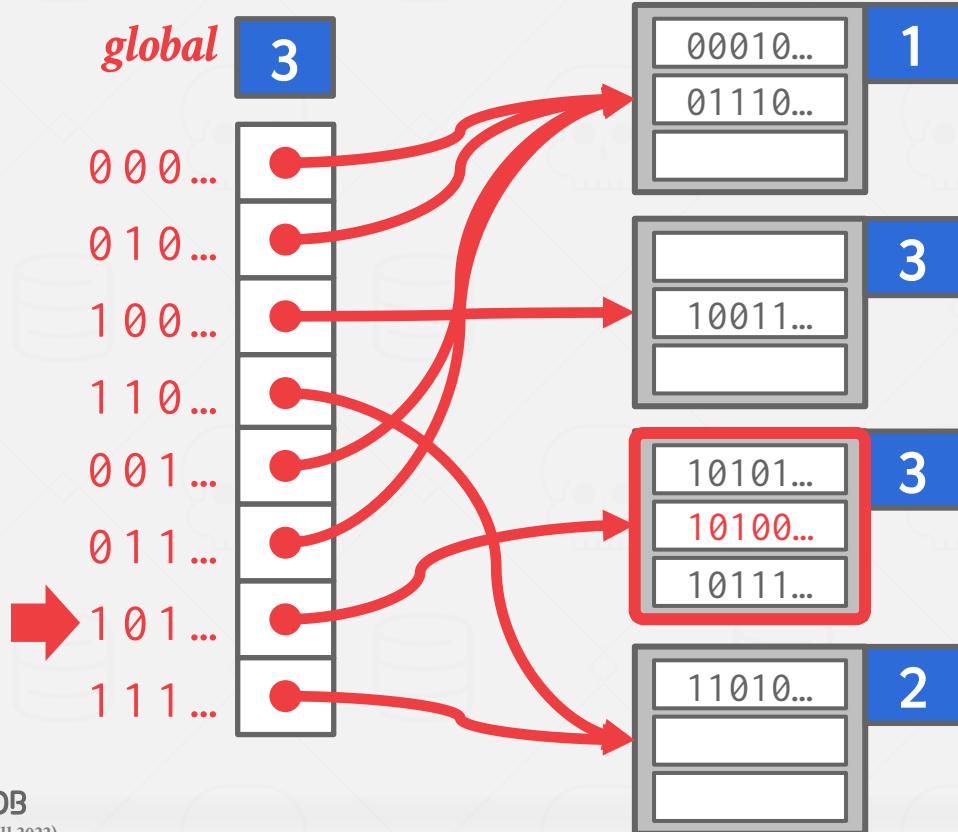
Put B

$\text{hash}(B) = 10111...$

Put C

$\text{hash}(C) = 10100...$

EXTENDIBLE HASHING



Get A

$\text{hash}(A) = 01110...$

Put B

$\text{hash}(B) = 10111...$

Put C

$\text{hash}(C) = \boxed{101}00...$

Explanation of how this

Mapping occurs ??

We use a hash table to maintain reference to the next bucket we want to split.

LINEAR HASHING

When any bucket overflows we split the bucket at that

The hash table maintains a pointer that tracks the location.
next bucket to split.

- When any bucket overflows, split the bucket at the pointer location.

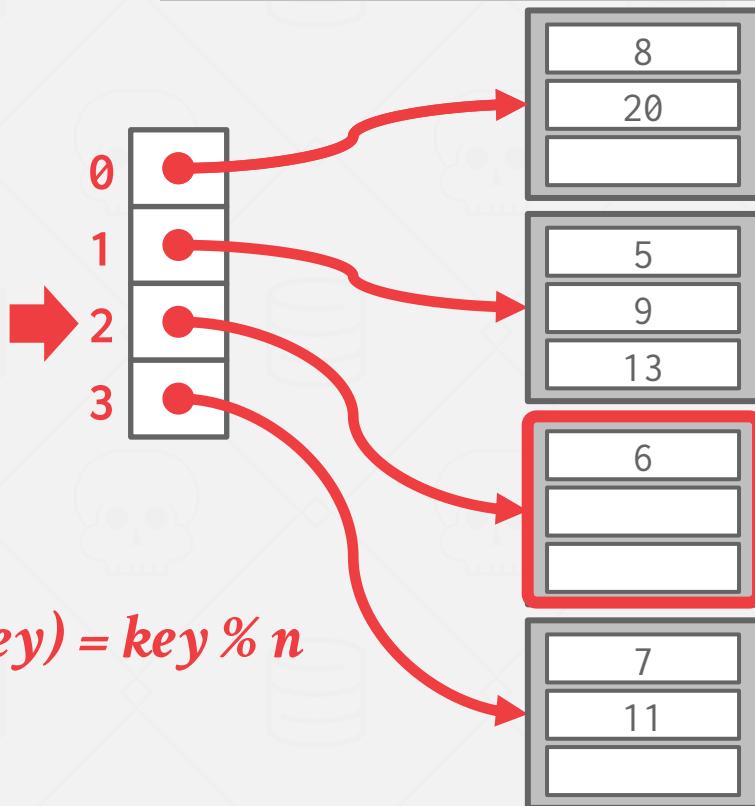
Use multiple hashes to find the right bucket for a given key.

Can use different overflow criterion:

- Space Utilization
- Average Length of Overflow Chains

LINEAR HASHING

*Split
Pointer*

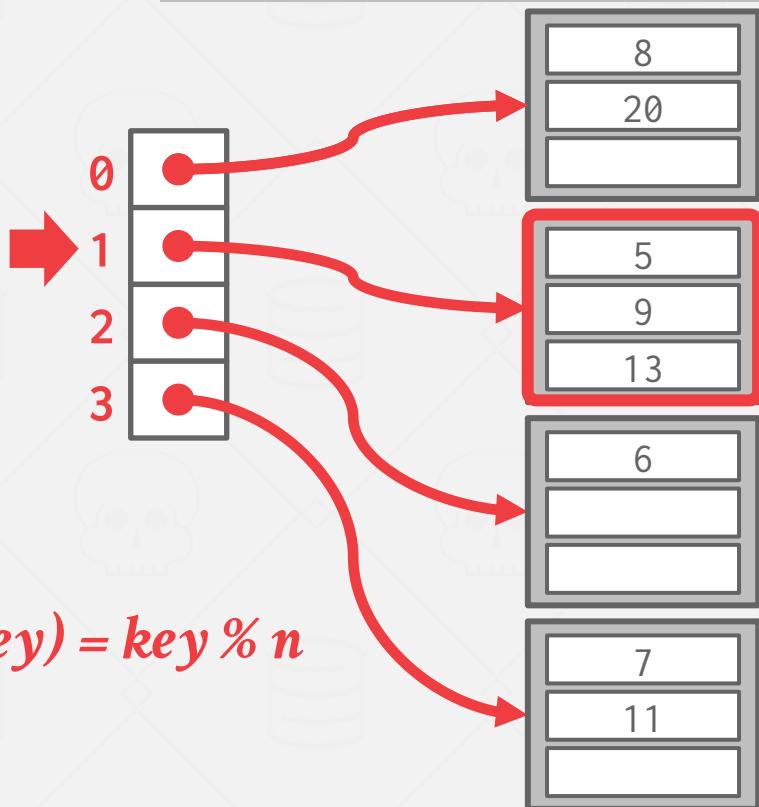


$$\text{hash}_1(\text{key}) = \text{key} \% n$$

Get 6
 $\text{hash}_1(6) = 6 \% 4 = 2$

LINEAR HASHING

*Split
Pointer*



$$\text{hash}_1(\text{key}) = \text{key} \% n$$

Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

$$\text{hash}_1(17) = 17 \% 4 = 1$$

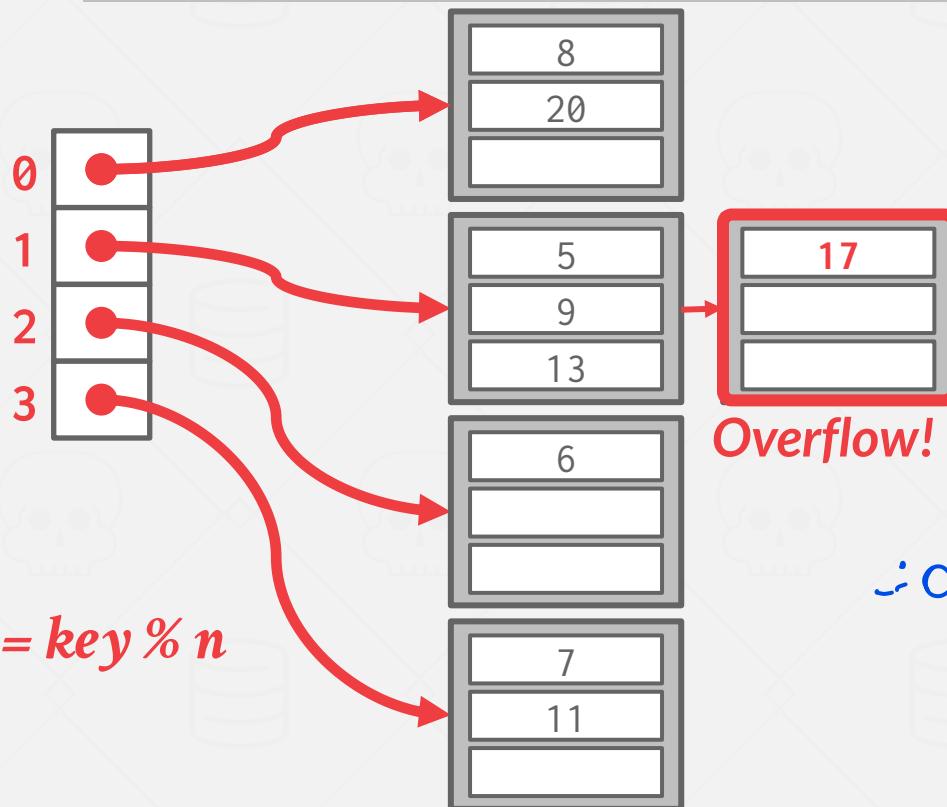
But in bucket 17 met there but

bucket overflow split

bucket 1 -

LINEAR HASHING

*Split
Pointer*



Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

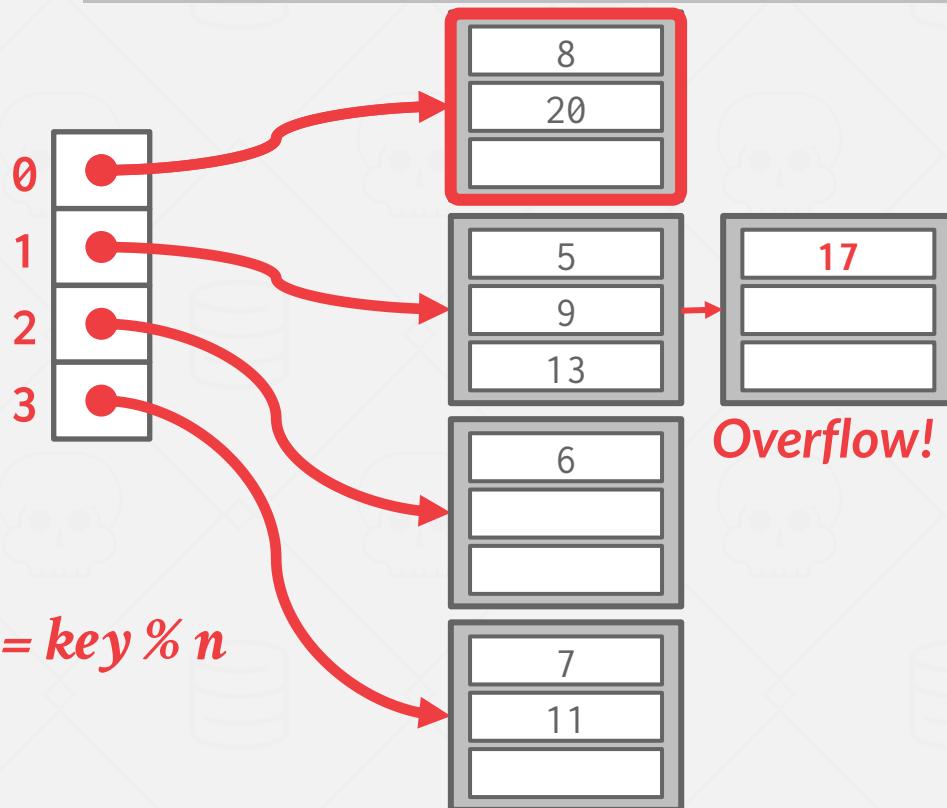
$$\text{hash}_1(17) = 17 \% 4 = 1$$

Overflow!

*∴ overflow 17 placed in
Split.*

LINEAR HASHING

*Split
Pointer*



Get 6

$$hash_1(6) = 6 \% 4 = 2$$

Put 17

$$hash_1(17) = 17 \% 4 = 1$$

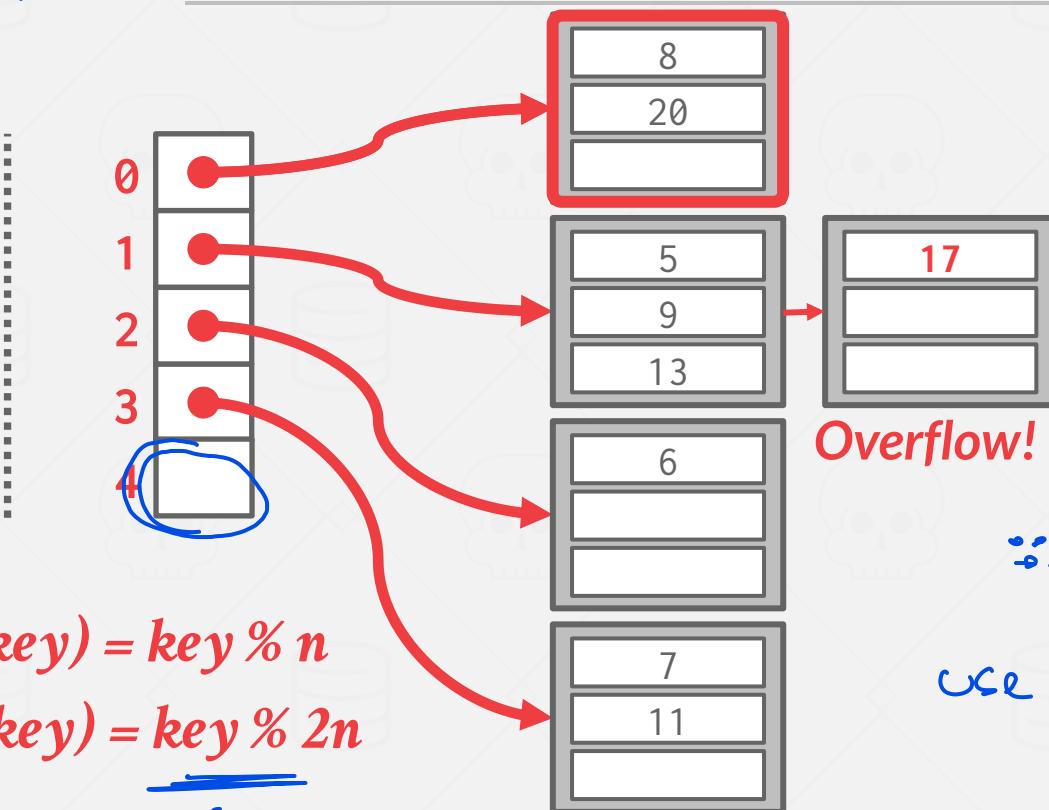
Overflow!

$$hash_1(key) = key \% n$$

LINEAR HASHING

*Split
Pointer*

$n \uparrow \text{by } 1$



$$\text{hash}_1(\text{key}) = \text{key \% } n$$

$$\text{hash}_2(\text{key}) = \underline{\text{key \% } 2n}$$

Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

$$\text{hash}_1(17) = 17 \% 4 = 1$$

Overflow!

∴ Split happened now we

use multiple hash functions.

↳ How did we design this hash 2 ?? for input key -

? Then based on hash2

We are rearranging

Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

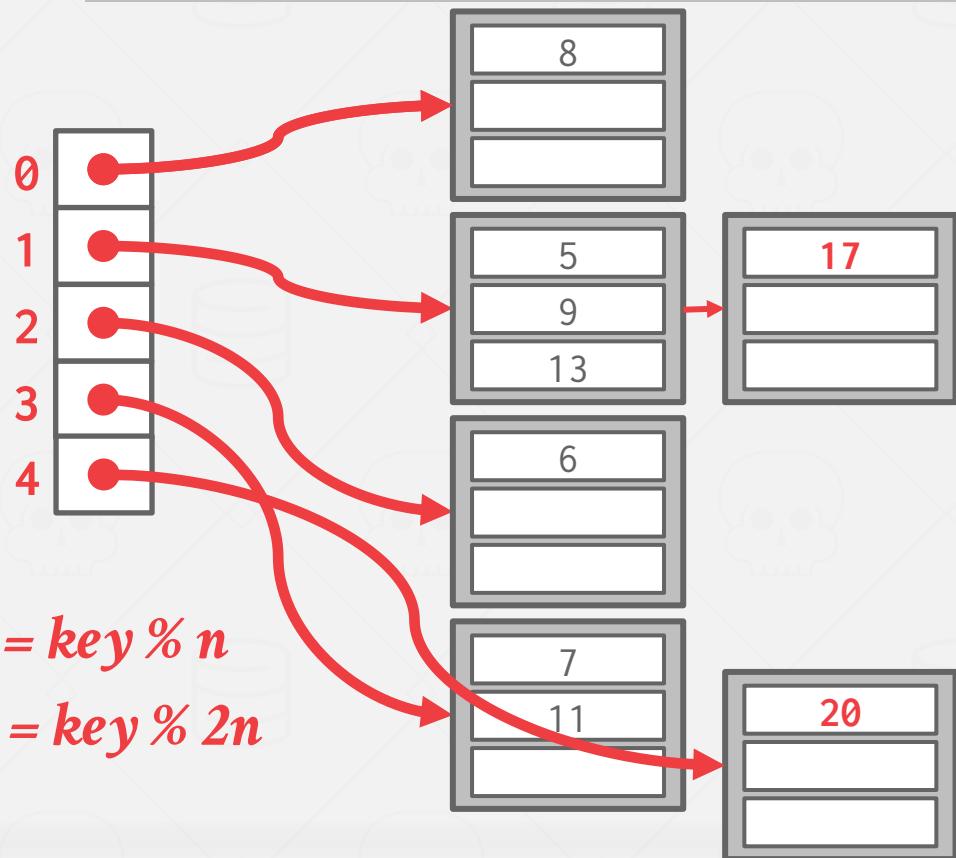
also,

Put 17

$$\text{hash}_1(17) = 17 \% 4 = 1$$

a previous key.

Split
Pointer

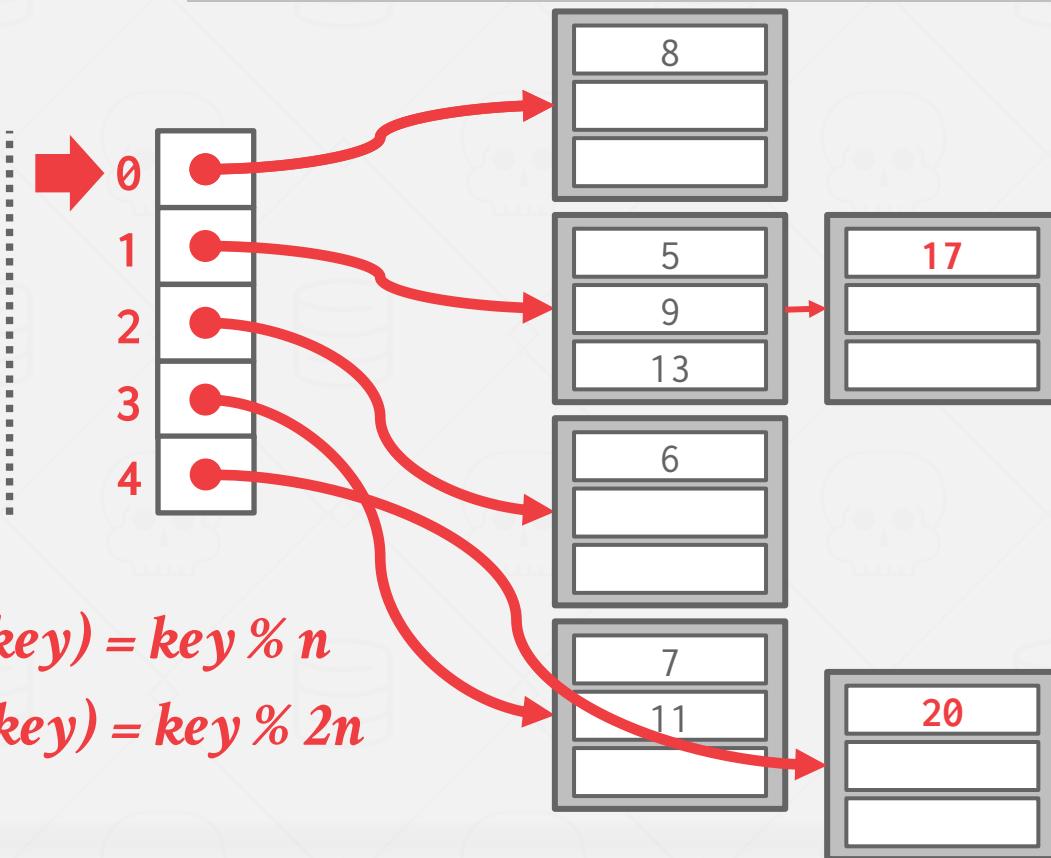


$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

LINEAR HASHING

*Split
Pointer*



$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

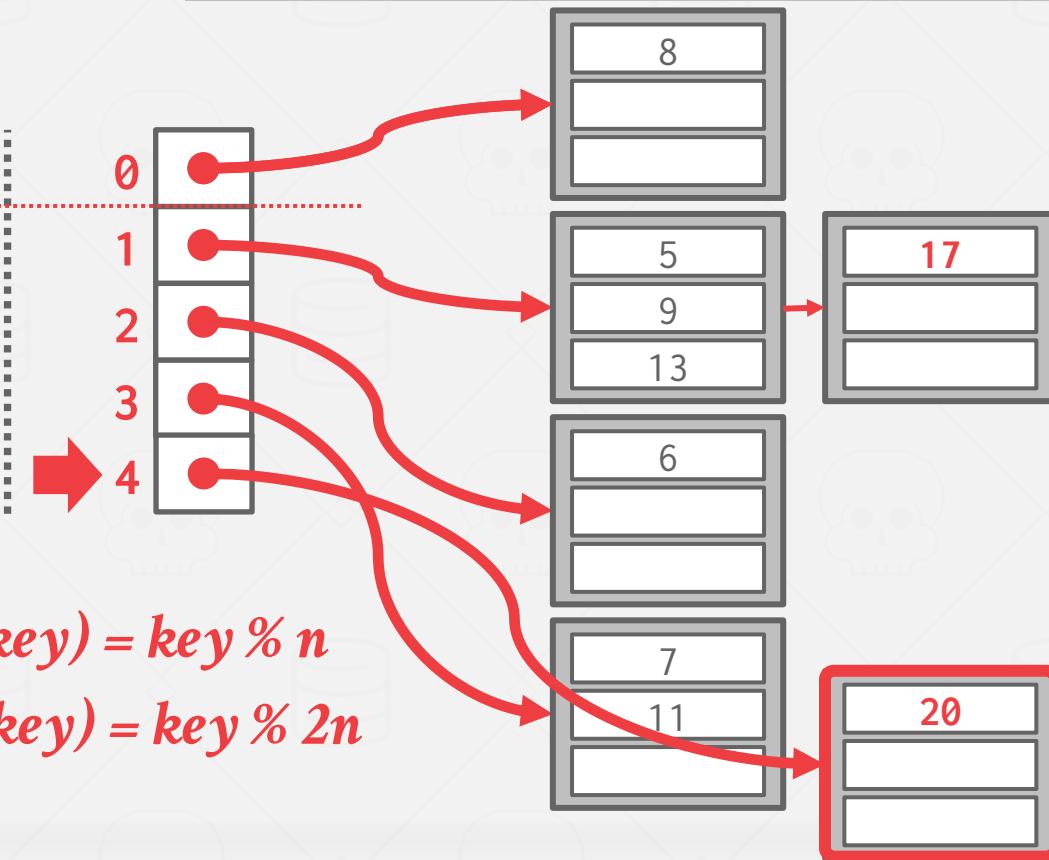
$$\text{hash}_1(17) = 17 \% 4 = 1$$

Get 20

$$\underline{\text{hash}_1(20) = 20 \% 4 = 0}$$

LINEAR HASHING

*Split
Pointer*



Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

$$\text{hash}_1(17) = 17 \% 4 = 1$$

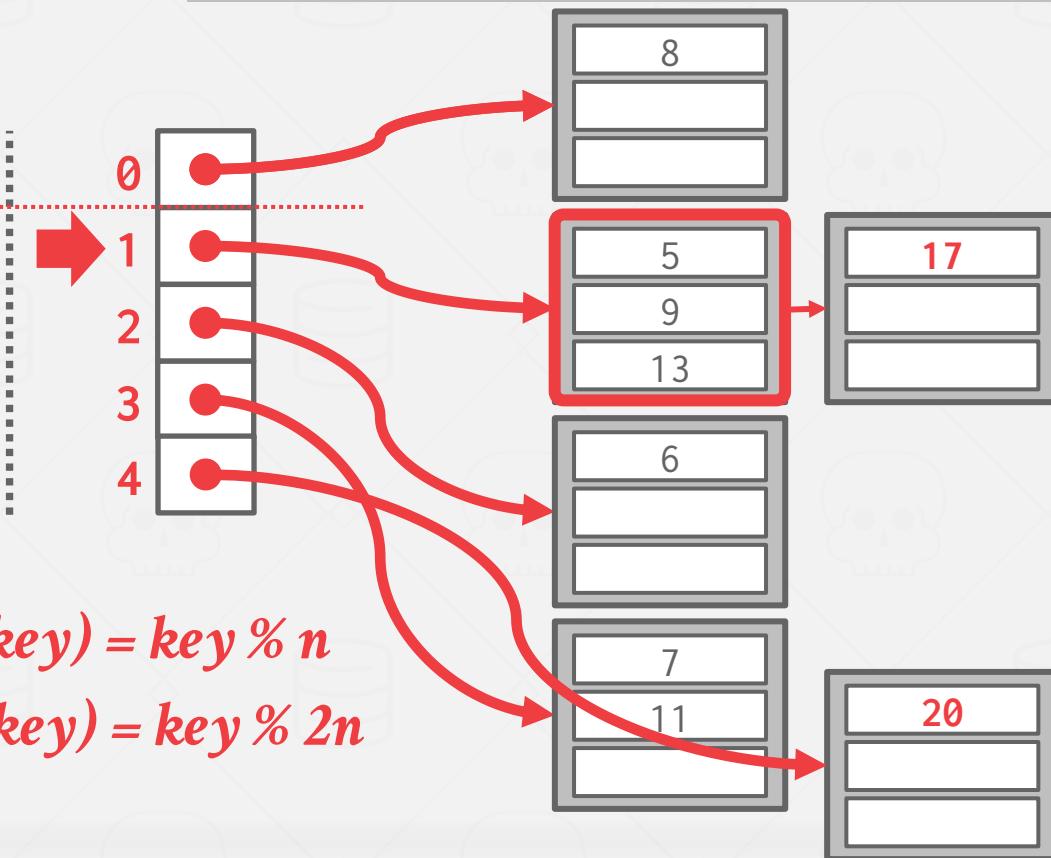
Get 20

$$\begin{aligned}\text{hash}_1(20) &= 20 \% 4 = 0 \\ \text{hash}_2(20) &= 20 \% 8 = 4\end{aligned}$$

2 different values -

LINEAR HASHING

*Split
Pointer*



$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

Get 6

$$\text{hash}_1(6) = 6 \% 4 = 2$$

Put 17

$$\text{hash}_1(17) = 17 \% 4 = 1$$

Get 20

$$\text{hash}_1(20) = 20 \% 4 = 0$$

$$\text{hash}_2(20) = 20 \% 8 = 4$$

Get 9

$$\text{hash}_1(9) = 9 \% 4 = 1$$

LINEAR HASHING

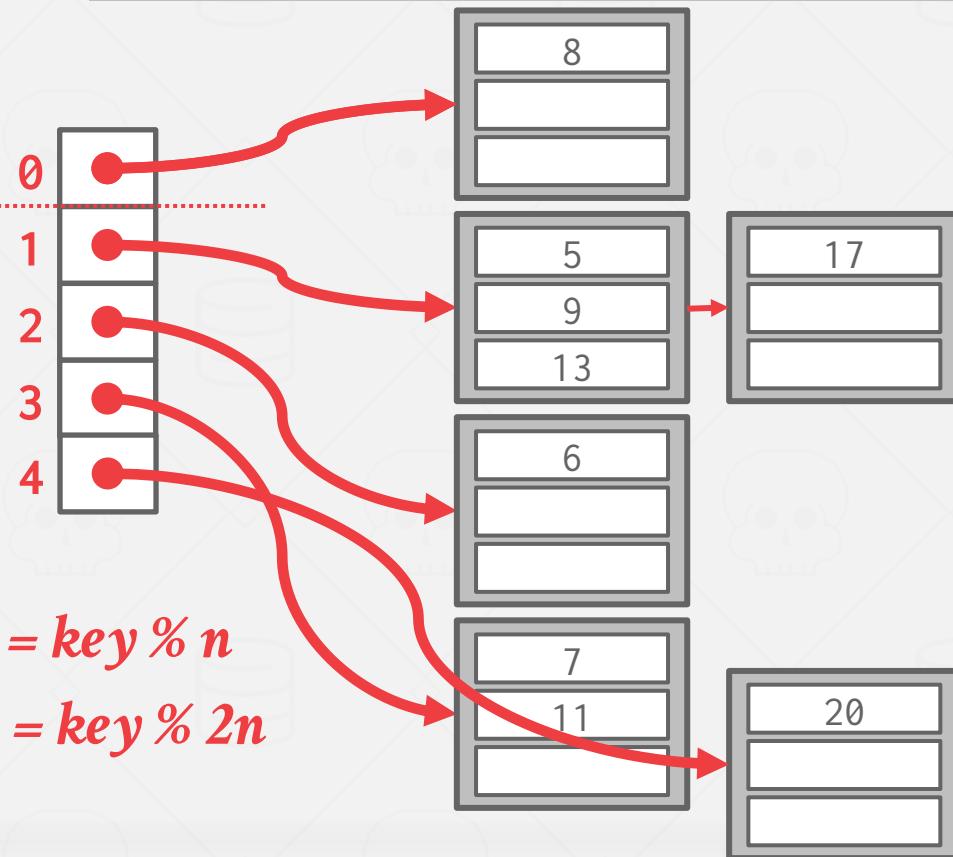
Splitting buckets based on the split pointer will eventually get to all overflowed buckets.

- When the pointer reaches the last slot, delete the first hash function and move back to beginning.

Did not understand Linear
Hashing ↗

LINEAR HASHING - DELETES

*Split
Pointer*

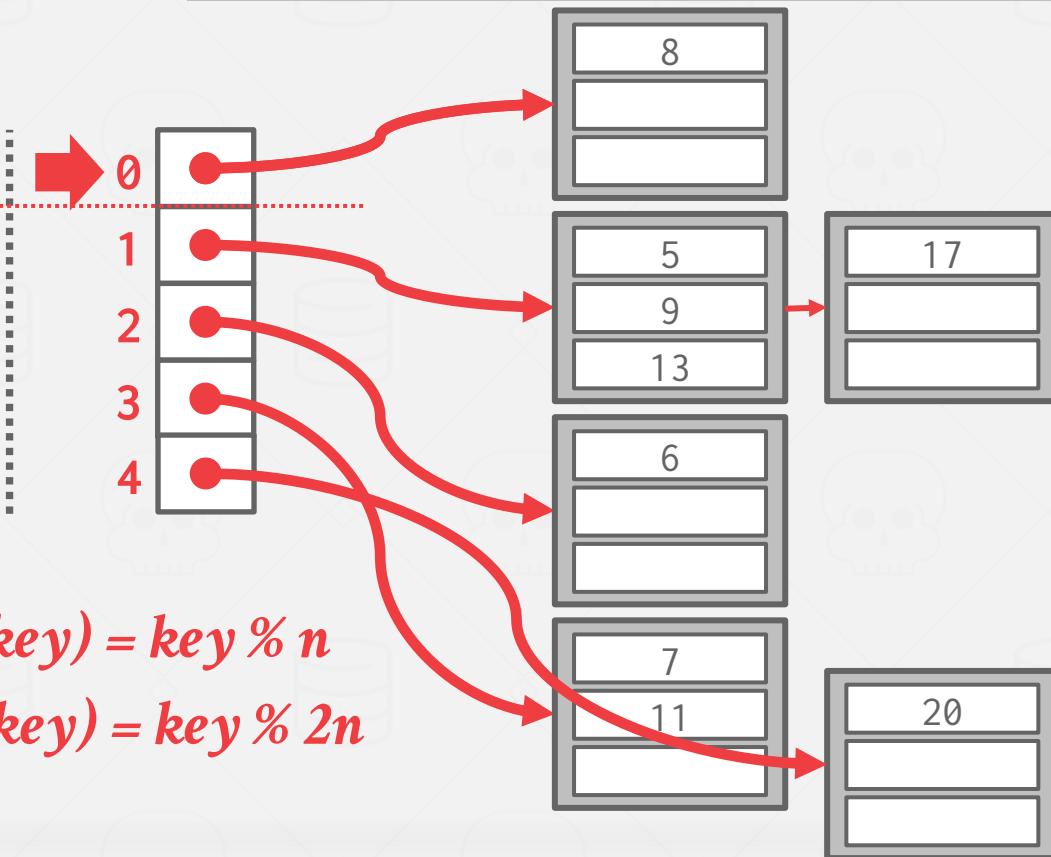


$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

LINEAR HASHING - DELETES

*Split
Pointer*



Delete 20

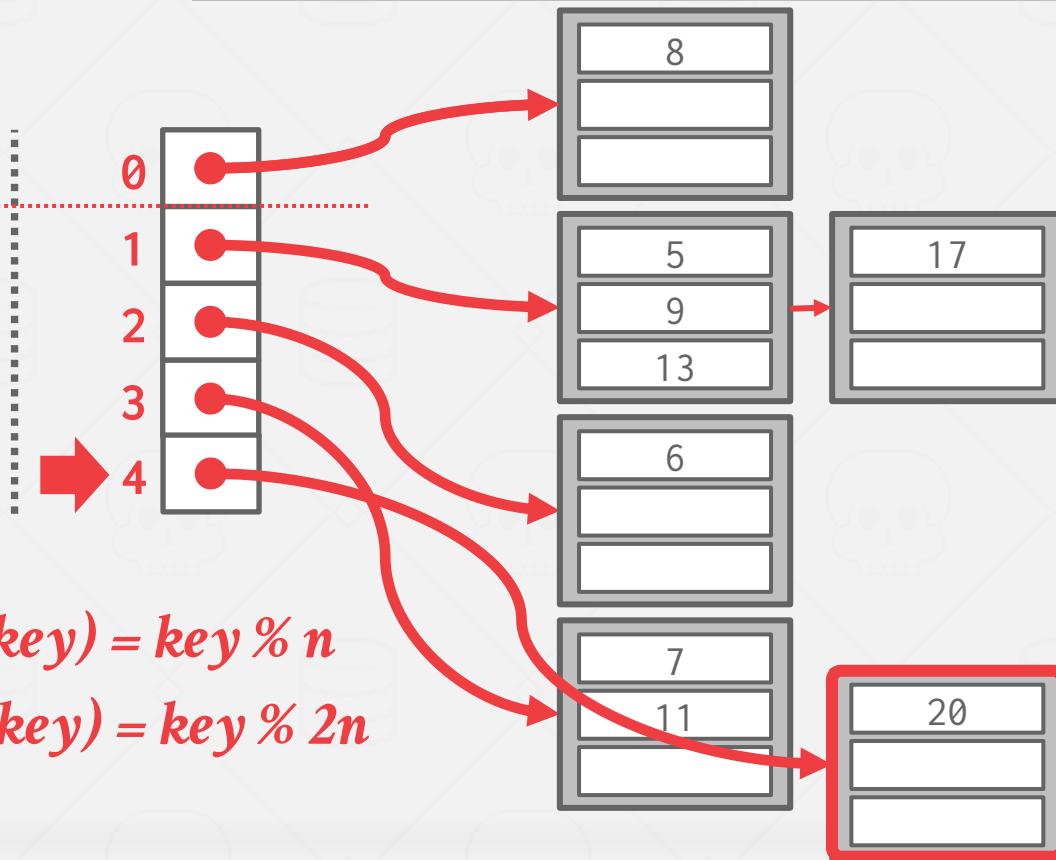
$$\text{hash}_1(20) = 20 \% 4 = 0$$

$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

LINEAR HASHING - DELETES

*Split
Pointer*



$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

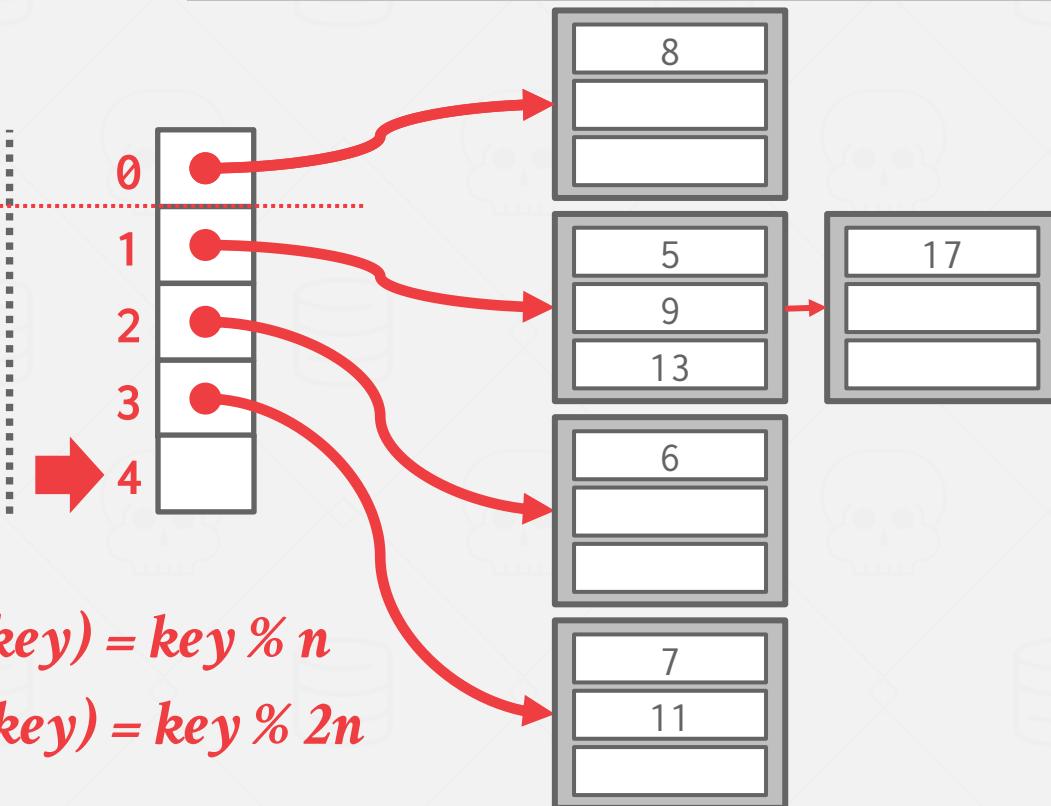
Delete 20

$$\text{hash}_1(20) = 20 \% 4 = 0$$

$$\text{hash}_2(20) = 20 \% 8 = 4$$

LINEAR HASHING - DELETES

*Split
Pointer*



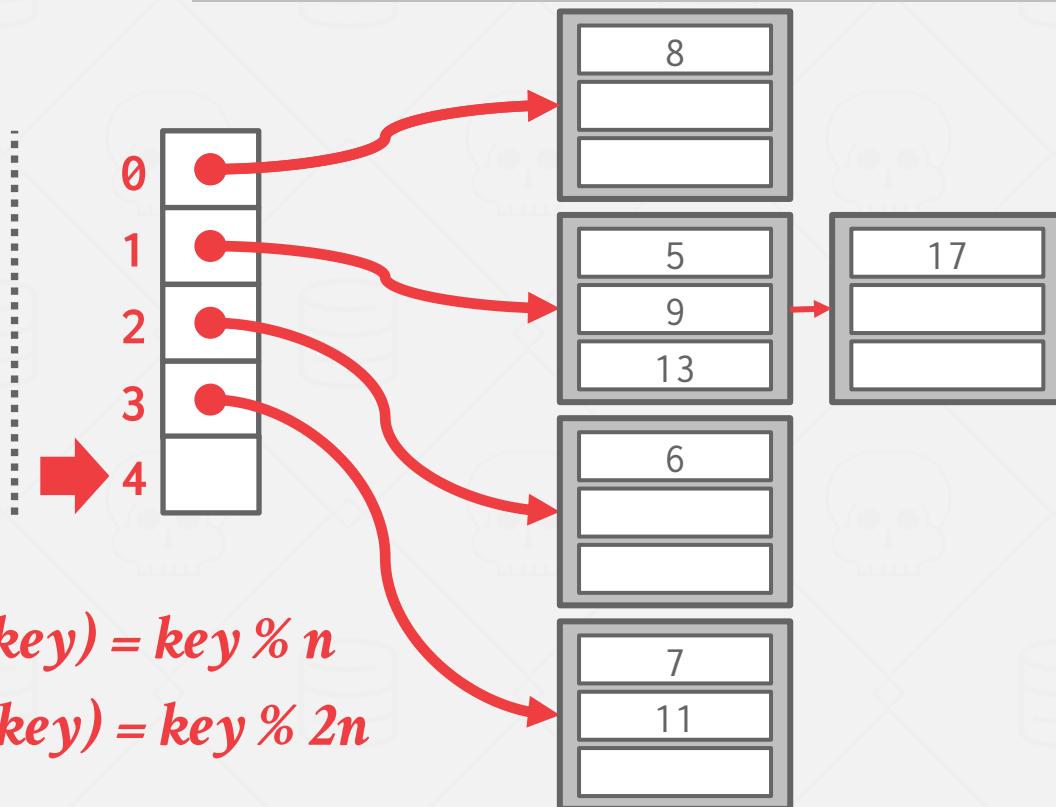
Delete 20

$$\begin{aligned} \text{hash}_1(20) &= 20 \% 4 = 0 \\ \text{hash}_2(20) &= 20 \% 8 = 4 \end{aligned}$$

$$\begin{aligned} \text{hash}_1(\text{key}) &= \text{key \% } n \\ \text{hash}_2(\text{key}) &= \text{key \% } 2n \end{aligned}$$

LINEAR HASHING - DELETES

*Split
Pointer*



$$\text{hash}_1(\text{key}) = \text{key} \% n$$

$$\text{hash}_2(\text{key}) = \text{key} \% 2n$$

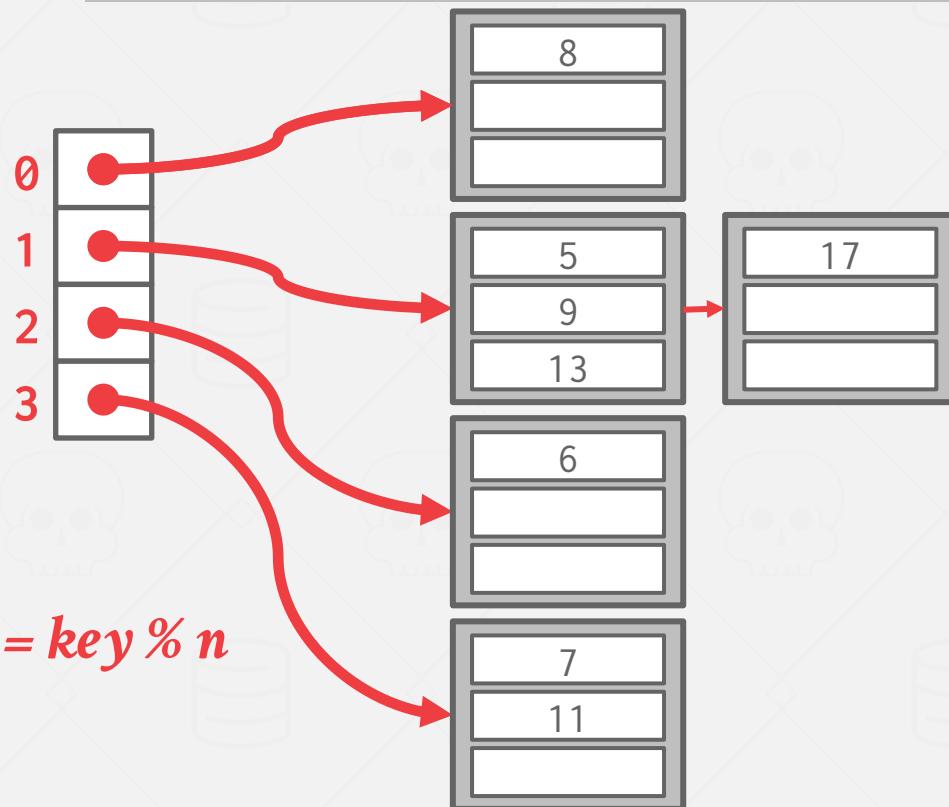
Delete 20

$$\text{hash}_1(20) = 20 \% 4 = 0$$

$$\text{hash}_2(20) = 20 \% 8 = 4$$

LINEAR HASHING - DELETES

*Split
Pointer*



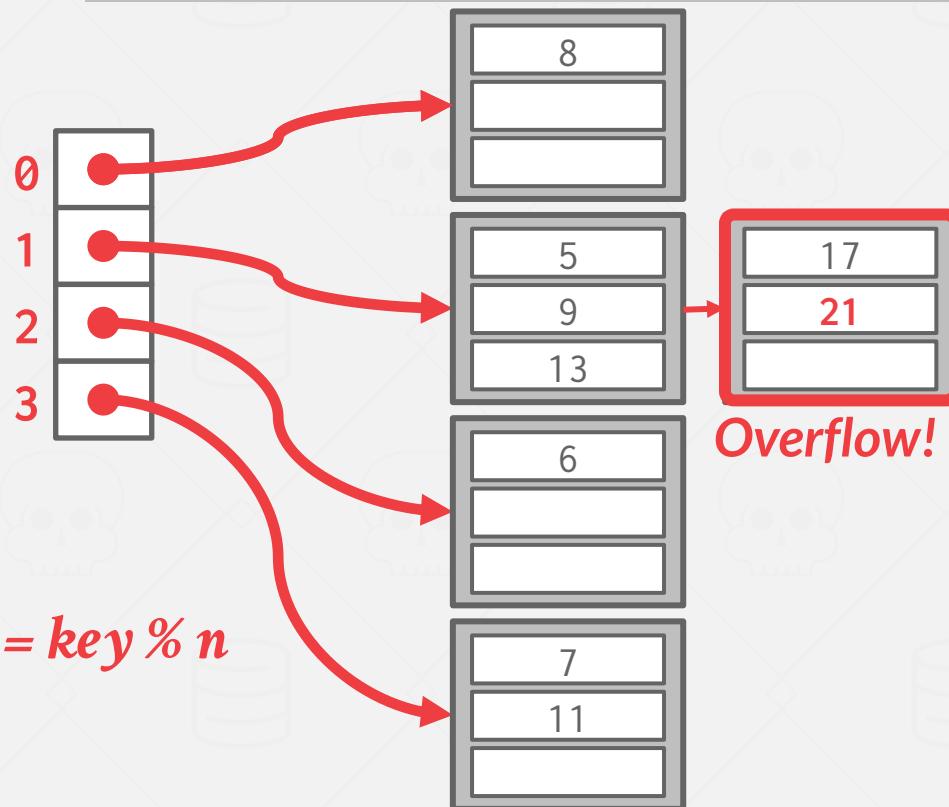
Delete 20

$$hash_1(20) = 20 \% 4 = 0$$

$$hash_2(20) = 20 \% 8 = 4$$

LINEAR HASHING - DELETES

*Split
Pointer*



Delete 20

$$hash_1(20) = 20 \% 4 = 0$$

$$hash_2(20) = 20 \% 8 = 4$$

Put 21

$$hash_1(21) = 21 \% 4 = 1$$

CONCLUSION

Fast data structures that support **O(1)** look-ups that are used all throughout DBMS internals.
→ Trade-off between speed and flexibility.

Hash tables are usually not what you want to use for a table index...

NEXT CLASS

B+Trees

→ aka "The Greatest Data Structure of All Time"