

Deep Learning

Pournami P N

Associate Professor

Department of Computer Science and Engineering

National Institute of Technology Calicut

NIT Campus Post, Kerala

pournamipn@nitc.ac.in

Introduction

- Inventors have long dreamed of creating machines that think - dates back to at least the time of ancient Greece.
- When programmable computers were first conceived - such machines might become intelligent.
- Today Artificial intelligence (AI) is a thriving field
 - intelligent software to automate routine labor
 - understand speech or images
 - make diagnoses in medicine
 - support basic scientific research

Challenges

- Early days of artificial intelligence
 - rapidly tackled and solved problems that are intellectually difficult for human beings but relatively straight-forward for computers
 - problems that can be described by a list of formal, mathematical rules.
 - *IBM's Deep Blue chess-playing system defeated world champion Garry Kasparov in 1997*
 - Chess can be completely described by a very brief list of completely formal rules, easily provided ahead of time by the programmer.
- **The true challenge to artificial intelligence**
- **Solving the tasks that are easy for people to perform but hard for people to describe formally**
 - problems that we solve intuitively, that feel automatic, like recognizing spoken words or faces in image

Challenges

- Abstract and formal tasks that are among the most difficult mental undertakings for a human being are among the easiest for a computer.
- Computers defeat even the best human chess player - but only recently matching some of the abilities of average human beings to recognize objects or speech.
- A person's everyday life requires an immense amount of knowledge about the world, which is subjective and intuitive, and therefore difficult to articulate in a formal way.
- Computers need to capture this same knowledge in order to behave in an intelligent way.
- *One of the key challenges in artificial intelligence is how to get this informal knowledge into a computer.*

Knowledge base Approach to AI

- Hard-code knowledge about the world in formal languages
- A computer can reason automatically about statements in these formal languages using logical inference rules.
 - Cyc (Lenat and Guha, 1989) failed to understand a story about a person named Fred shaving in the morning (Linde, 1992)

Solution

- Allow computers to learn from experience
- Understand the world in terms of a hierarchy of concepts, with each concept defined through its relation to simpler concepts.
- This approach avoids the need for human operators to formally specify all the knowledge that the computer needs.
- The hierarchy of concepts enables the computer to learn complicated concepts by building them out of simpler ones.
- This is done by extracting patterns from raw data - known as ***MACHINE LEARNING*** (subjective decisions by computer programs)

Learning Algorithm

Learning

What is the Learning Problem?

Learning = Improving with experience at some task
Improve over task T ,
with respect to performance measure P ,
based on experience E.

E.g., Learn to play checkers

T : Play checkers

P : % of games won in world tournament

E: opportunity to play against self

Lecture slides for textbook Machine Learning, T. Mitchell, McGraw Hill, 1997

- A simple machine learning algorithm called logistic regression can determine whether to recommend cesarean delivery ([Mor-Yosef et al., 1990](#)).
- A simple machine learning algorithm called naive Bayes can separate legitimate e-mail from spam e-mail

Representation of Data

- The **performance of the machine learning algorithms** depends on the **representation of the data** they are given.
 - When logistic regression is used to recommend cesarean delivery, the AI system does not examine the patient directly.
 - Instead, the doctor tells the system several pieces of relevant information, such as the presence or absence of a uterine scar.
 - Each piece of information included in the representation of the patient is known as a **feature**.
 - Logistic regression learns how each of these features of the patient correlates with various outcomes

Representation of Data

- The choice of representation has an enormous effect on the performance of machine learning algorithms

Representations Matter




Figure 1.1

(Goodfellow 2016)

- Many artificial intelligence tasks
 - designing the right set of features to extract for that task, providing these features to a simple machine learning algorithm.
 - Speaker identification from sound - an estimate of the size of the speaker's vocal tract.
 - This feature gives a strong clue - the speaker is a man, woman, or child.

Representation of Data

- **What features should be extracted?**

- Write a program to detect cars in photographs.

- **Solution - Representation Learning**

- Use machine learning to discover not only the mapping from representation to output but also the representation itself.
- Enable AI systems to rapidly adapt to new tasks, with minimal human intervention.

Factors of Variation

- When designing features or algorithms for learning features, our goal is usually to *separate the factors of variation* that explain the observed data.
- Many of the factors of variation influence every single piece of data we are able to observe.
- Many of these factors of variation, such as a speaker's accent, can be identified only using sophisticated, nearly human-level understanding of the data.
- **Deep learning solves this central problem in representation learning by introducing representations that are expressed in terms of other, simpler representations.**
 - **Multi-Layer Perceptron (MLP)**

Depth: Repeated Composition

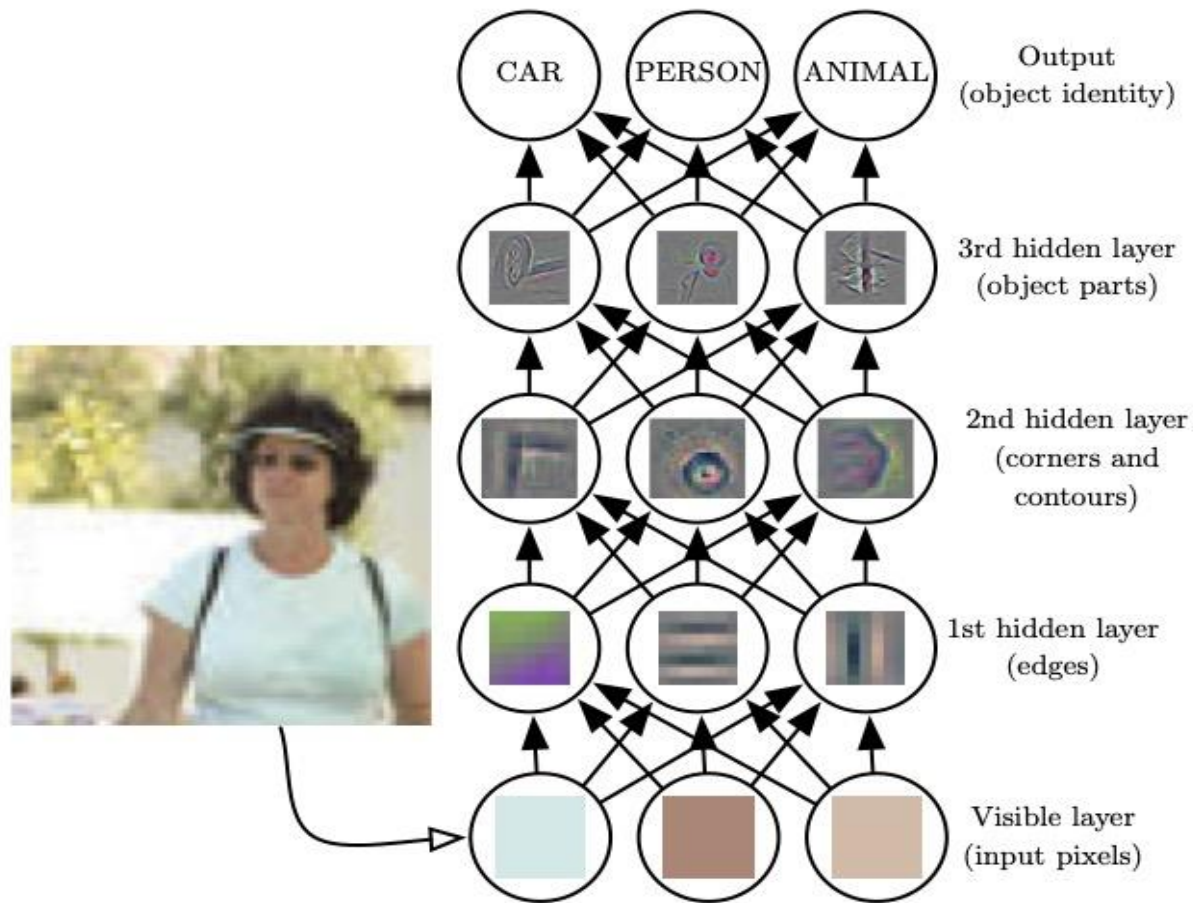


Figure 1.2

(Goodfellow 2016)

Zeiler and Fergus (2014)

Deep learning is an approach to AI

Machine Learning and AI

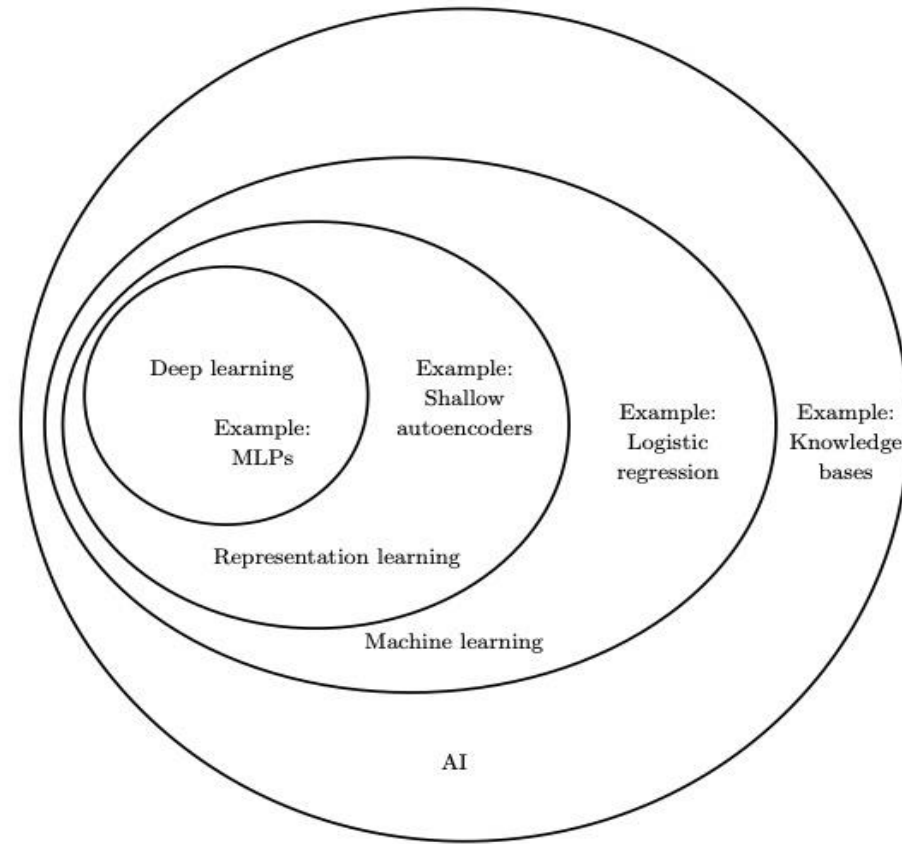


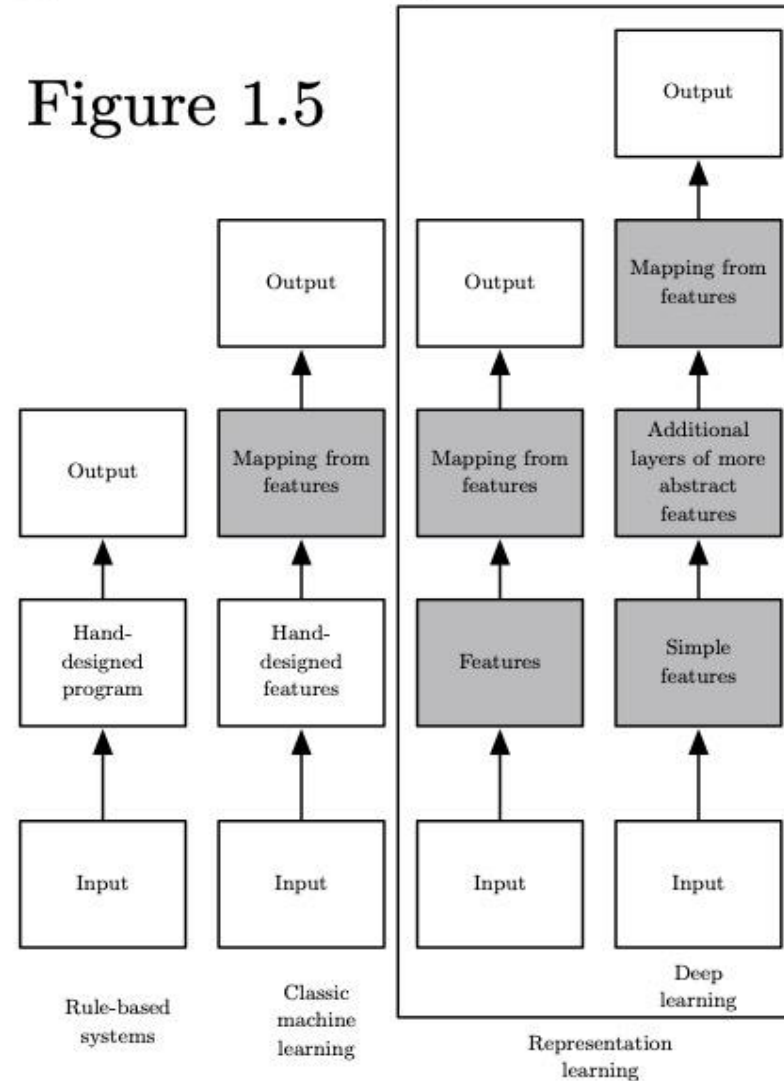
Figure 1.4

(Goodfellow 2016)

Learning Multiple Components

Parts
of an
AI
system

Figure 1.5



Conclusion

- Machine learning is the only viable approach to build AI systems that can operate in complicated real-world environments.
- Deep learning is a particular kind of machine learning
 - achieves great power and flexibility by representing the world as a nested hierarchy of concepts, with each concept defined in relation to simpler concepts
- Deep Learning makes it possible for more abstract representations computed in terms of less abstract ones.